

# *Introduction to Data Management*

## *Lecture #4* *E-R Model, Still Going...*



Instructor: Mike Carey  
[mjcarey@ics.uci.edu](mailto:mjcarey@ics.uci.edu)

# Today's Reminders




- ❖ Continue to follow the course wiki page
  - <http://www.ics.uci.edu/~cs122a/>
- ❖ Also follow (and live by) the Piazza page
  - <https://piazza.com/uci/spring2019/cs122a/home>
- ❖ The first HW assignment is underway
  - Conceptual (E-R) database design for **PHLOG**
- ❖ This week's quiz will be on E-R modeling
  - Also bring any lingering HW #1 related questions

# *More on Assignment 1*




- ❖ Due date: April 12th @ 5PM
- ❖ Late submission: April 13th @ 5PM with 20% point deduction
- ❖ Use software to draw your E-R diagram in the provided template and follow the instructions carefully (including Entity placement!)
- ❖ Submit your assignment to Gradescope
- ❖ Everyone should have received Gradescope (if not, join using the following code: MZYZW6)


# Add'l Advanced ER Features (Review)




- ❖ Multi-valued (vs. single-valued) attributes




- ❖ Optional attributes




- ❖ Composite (vs. atomic) attributes



- ❖ Derived (vs. base/stored) attributes






# Conceptual Design Using the ER Model

## ❖ Design choices:

- Should a given concept be modeled as an entity or as an attribute?
- Should a given concept be modeled as an entity or as a relationship?
- Characterizing relationships: Binary or ternary? Aggregation? ...

## ❖ Constraints in the ER Model:


- A lot of data semantics<sup>语义</sup> can (and should) be captured
- But, not all constraints can be captured by ER diagrams. (Ex: Department heads from earlier...!)




# Advanced Attribute Considerations

- ❖ Should *address* be an attribute of Employees or an entity (connected to Employees by a relationship)?
- ❖ Depends on how we want to use addresses, on the data semantics, and what model features we use:
  - If we have several addresses per employee, *address* would be a multivalued attribute or an entity if we stick only to basic E-R concepts (as they can't be set-valued w/o advanced modeling).
  - If address structure (city, street, etc.) is important, e.g., to query for employees in a given city, *address* should be modeled as a **composite attribute or an entity** (as attribute values are *atomic* otherwise) – i.e., it shouldn't just be an address string.
  - If the address itself is logically separate (e.g., representing a property that's located there) and refer-able, it's rightly an entity in any case!

# Attribute Considerations (Cont'd.)



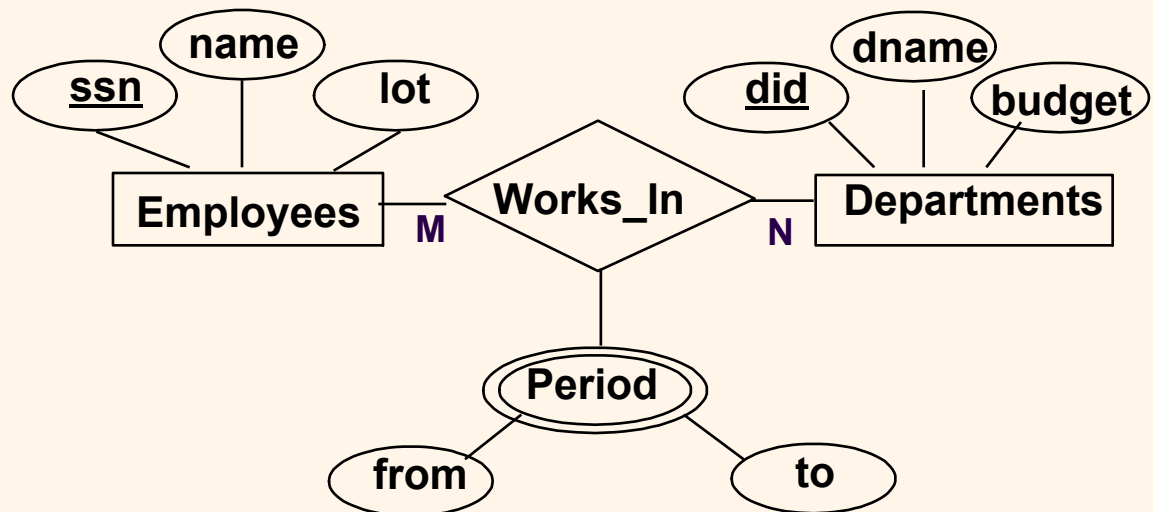
- ❖ Works\_In here does not allow an employee to work in a department for two or more periods. (Q: Why...?)



因为他只有一个from一个to


- ❖ Similar to the problem of wanting to record several addresses for an employee: We want to record *several values of the descriptive attributes for each instance of this "relationship"*.

Could be accomplished by having a *multivalued* relationship attribute:



relationship直接的动作会产生副产品，就讲这个副产品为动作的attribute

# Attribute Considerations (Cont'd.<sup>2</sup>)

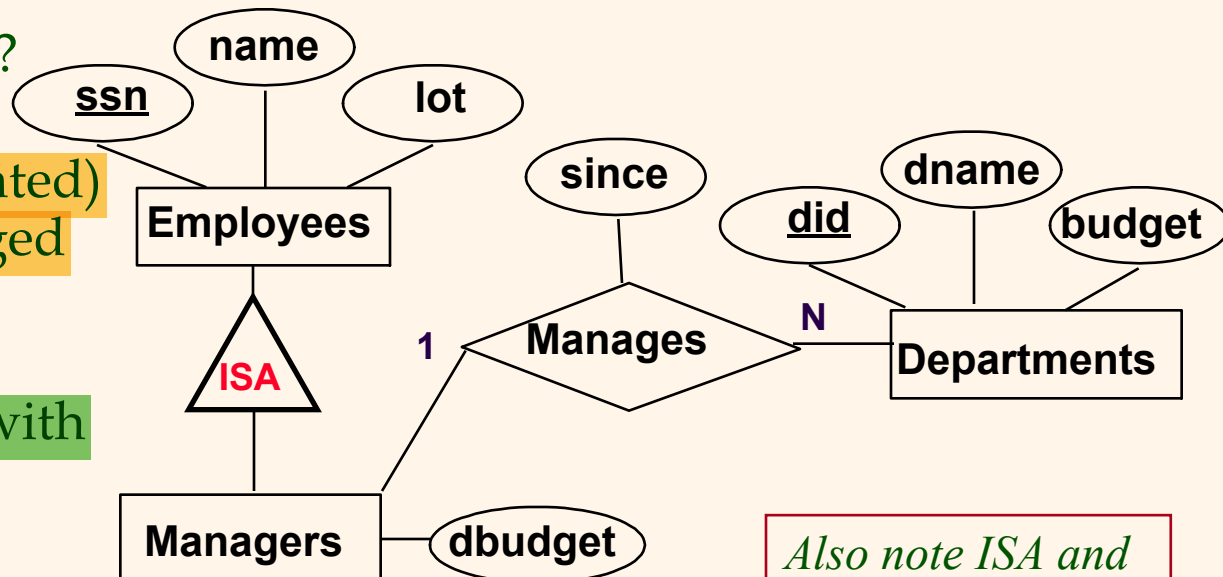



这个dbudget是某个employees对某个department的dbudget，如果这个员工还有其他department的话，也有其他department的dbudget，所以这里的dbudget是指某个员工对于某个单独的一个的department的钱

❖ ER diagram on the right is okay if a manager gets a separate discretionary budget (dbudget) per dept.

❖ What if a manager gets a discretionary budget that covers *all* managed depts?


- **Redundancy:** dbudget could be stored (repeated) with each dept managed by the manager.
- **Misleading:** Suggests dbudget is associated with department-mgr combination.



这个是manager的全部的钱可以来自单个或多个的department的钱

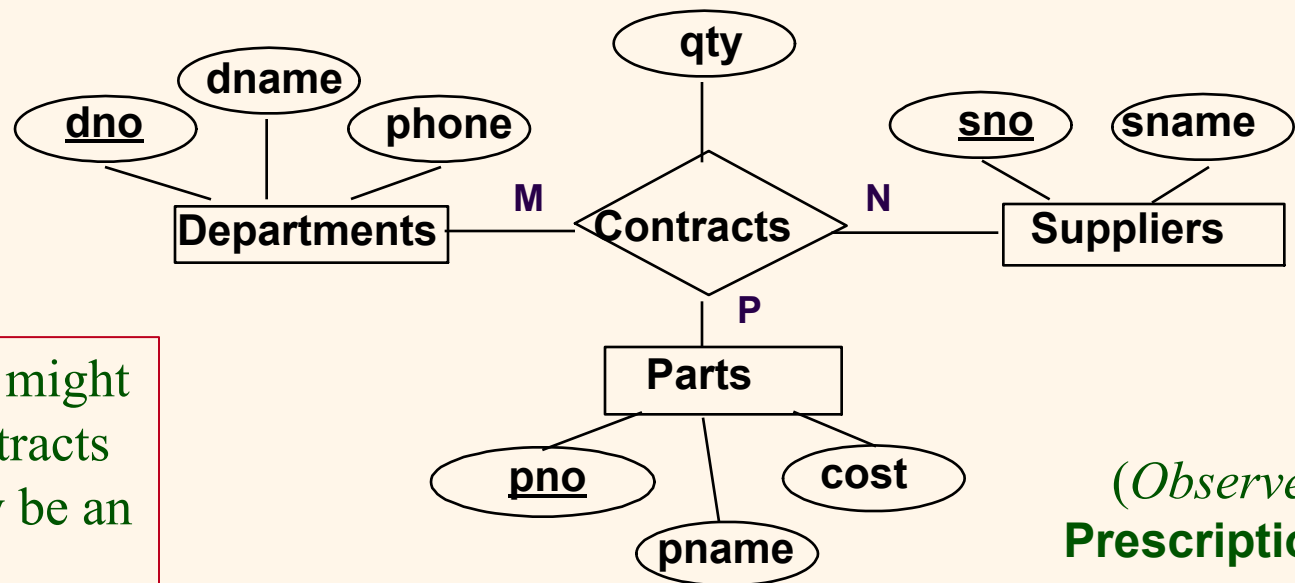
*Also note ISA and the relationship...*






# Binary vs. Ternary Relationships

- ❖ An example where ternary is needed: a ternary relation **Contracts** relates the entity sets **Parts**, **Departments** and **Suppliers**, and has descriptive attribute *qty*:



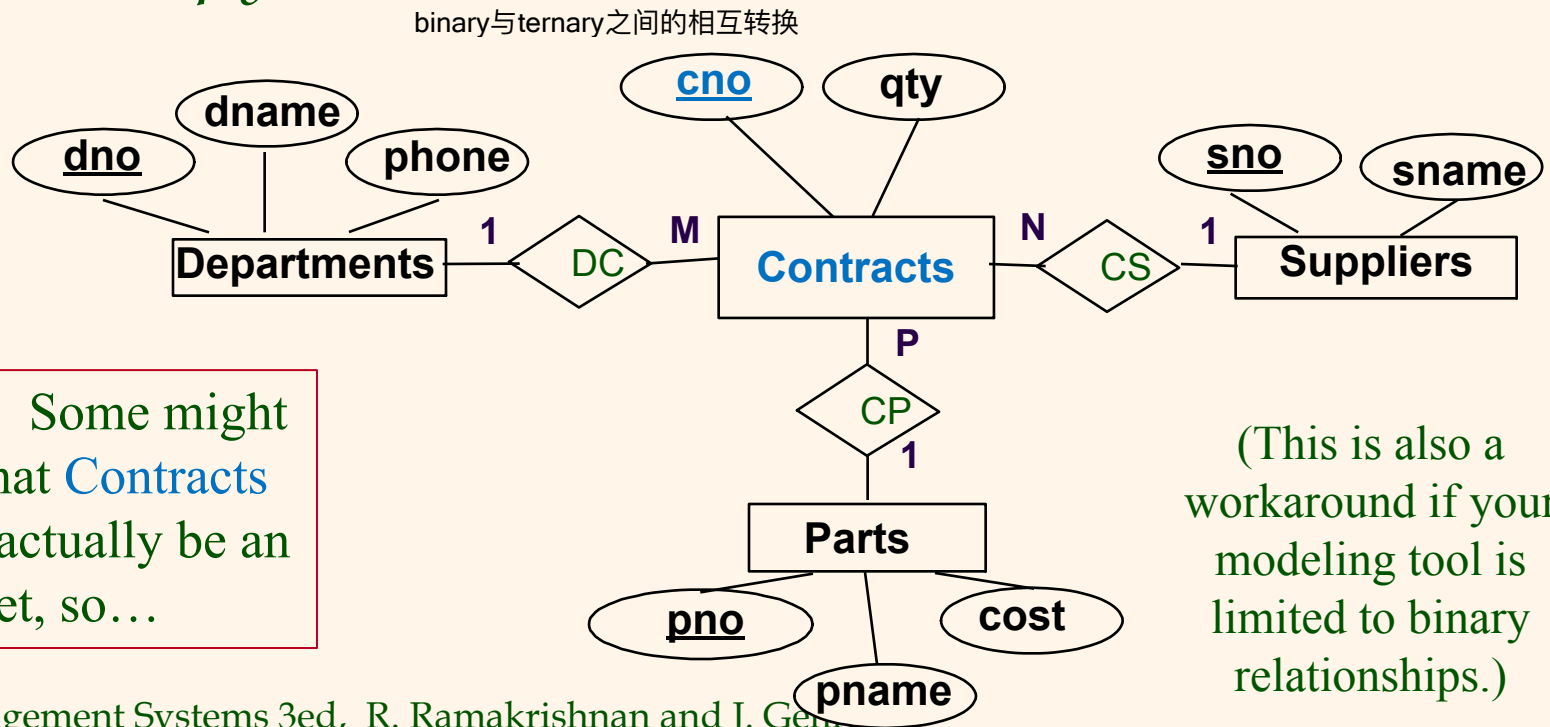
**NOTE:** Some might argue that Contracts should actually be an entity set...


(Observe:  
**Prescriptions**  
was similar)



## Binary vs. Ternary Relationships (Cont'd.)

- ❖ What the entity set perspective would lead to: an entity set **Contracts** related to the entity sets **Parts**, **Departments** and **Suppliers**, with the descriptive attribute *qty*:






# *Datatase Design Process (Flow)*

- ❖ Requirements Gathering (interviews)
- ❖ Conceptual Design (using ER)
- ❖ Platform Choice (which DBMS?)
- ❖ Logical Design (for target data model)
- ❖ Physical Design (for target DBMS, workload)
- ❖ Implement (and test, of course ☺)


database 设计的步骤

(Expect backtracking, iteration, and also incremental adjustments – and, we will actually be giving you a bit of practice with that last one in the next few HW assignments...! ☺)



# *Summary of Conceptual Design*

- ❖ *Conceptual design follows requirements analysis*
  - Yields a high-level description of data to be stored
- ❖ *ER model popular for conceptual design*
  - Constructs are expressive, close to the way people think about their applications.
- ❖ Basic constructs: *entities, relationships, and attributes* (of entities and relationships).
- ❖ Additionally: *weak entities, ISA hierarchies, aggregation, and multi-valued, composite and/or derived attributes.*
- ❖ Note: Many variations on the ER model (and many notations in use as well) – and also, UML...




# Summary of ER (Cont'd.)

完整性约束


- ❖ Several kinds of integrity constraints can be expressed in the ER model: *cardinality constraints*, *participation constraints*, also *overlap/covering constraints* for ISA hierarchies. Some *foreign key constraints* are also implicit in the definition of a relationship set (more about key constraints will be coming soon).
  - Some constraints (notably, *functional dependencies*) cannot be expressed in the ER model.
  - Constraints play an important role in determining the best database design for an enterprise.

企业




## Summary of ER (Cont'd.)

- ❖ ER design is *subjective*. There are often many ways to model a given scenario! Analyzing alternatives can be tricky, especially for a large enterprise. Common choices include:
  - Entity vs. attribute, entity vs. relationship, binary or n-ary relationship, whether or not to use an ISA hierarchy, and whether or not to use aggregation.
- ❖ Ensuring good database design: The resulting relational schema should be analyzed and refined further. For this, FD information and normalization techniques are especially useful (coming soon).



# Relational Database: Definitions

- ❖ Relational database: a set of *relations*
- ❖ Relation: consists of 2 parts:
  - *Instance*: a table, with rows and columns.  
#Rows = cardinality, #fields = degree or arity.
  - *Schema*: specifies name of relation, plus name and type of each column.
    - E.g. Students(*sid*: string, *name*: string, *login*: string, *age*: integer, *gpa*: real).
- ❖ Can think of a relation as a *set* of rows or *tuples* (i.e., all rows are distinct) in the pure relational model (*vs.* reality of SQL 😊)



## *Example Instance of Students Relation*


sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

❖ Cardinality = 3, degree = 5, all rows distinct

行数是 Cardinality = 3, 就是这个图总共能有多少种不同


degree是列, 表示有多少个attribute






# Relational Query Languages

- ❖ A major strength of the relational model: supports simple, powerful *querying* of data.
- ❖ Queries can be written intuitively, and the DBMS is responsible for efficient evaluation.
  - The key: precise (and set-based) semantics for relational queries. 关系查询的精确（和基于集合）语义
  - Allows the optimizer to extensively re-order operations, and still ensure that the answer does not change.



# *The SQL Query Language (Preview)*

- ❖ Developed by IBM (System R) in the 1970s
- ❖ Need for a standard, since it is used by many vendors (Oracle, IBM, Microsoft, ...)
- ❖ ANSI/ISO Standards:
  - SQL-86
  - SQL-89 (minor revision)
  - SQL-92 (major revision, very widely supported)
  - SQL-99 (major extensions, current standard)



# *The SQL Query Language (Preview)*

- ❖ To find all 18 year old students, we can write:

```
SELECT *  
FROM Students S  
WHERE S.age=18
```


sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@ee	18	3.2

select是指就找到什么，就把那一整列列出来，相符的都列出来

- To find just names and logins, replace the first line:

```
SELECT S.name, S.login
```

如果只要名字和login的话，那就可以是select S.name和S.login



# Querying Multiple Relations

- ❖ What does the following query compute?

```
SELECT S.name, E.cid  
FROM Students S, Enrolled E  
WHERE S.sid=E.sid AND E.grade="A"
```

Given the following instances of Students and Enrolled:

这个是 Student S

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@ecs	18	3.2
53650	Smith	smith@math	19	3.8


这个是Enrolled E

sid	cid	grade
53831	Carnatic101	C
53831	Reggae203	B
53650	Topology112	A
53666	History105	B

b不符合

We will get:

S.name	E.cid
Smith	Topology112




# Creating Relations in SQL

这样的可以create一个table

- ❖ Create the Students relation. Observe that the type (*domain*) of each field is specified and enforced by the DBMS whenever tuples are added or modified.
- ❖ As another example, the Enrolled table holds information about courses that students take.

```
CREATE TABLE Students  
(sid CHAR(20),  
 name CHAR(20),  
 login CHAR(10),  
 age INTEGER,  
 gpa REAL)
```

```
CREATE TABLE Enrolled  
(sid CHAR(20),  
 cid CHAR(20),  
 grade CHAR(2))
```



# *Destroying and Altering Relations*

## **DROP TABLE** Students

drop table 用来delete


- ❖ Destroys the relation Students. The schema information *and* the tuples are deleted.

## **ALTER TABLE** Students

### **ADD COLUMN** firstYear integer

加一个alter table 可以往后面加一个column (列) , 然后值都为null

- ❖ The schema of Students is altered by adding a new field; every tuple in the current instance is extended with a *null* value in the new field.



# *Adding and Deleting Tuples*

- ❖ Can insert a single tuple using:

```
INSERT INTO Students (sid, name, login, age, gpa)  
VALUES (53688, 'Smith', 'smith@ee', 18, 3.2)
```


可以增加一行数据

- ❖ Can delete all tuples satisfying some condition (e.g., name = Smith):

可以减去所有tuple满足以下关系

```
DELETE  
FROM Students S  
WHERE S.name = 'Smith'
```

- *Powerful variants of these commands are available; more later!*



# Integrity Constraints (ICs)

- ❖ **IC:** condition that must be true for *any* instance of the database; e.g., *domain constraints*.
  - ICs are specified when schema is defined.
  - ICs are checked when relations are modified.
- ❖ A *legal* instance of a relation is one that satisfies all specified ICs.
  - DBMS should not allow illegal instances.
- ❖ If the DBMS checks ICs, stored data is more faithful to real-world meaning.
  - Avoids data entry errors (centrally), too!