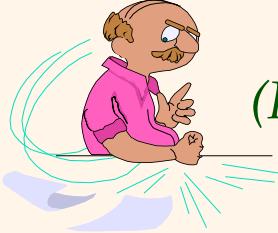


Introduction to Data Management

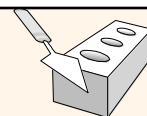


Lecture #12 (Relational Languages II)

Instructor: Mike Carey
mjcarey@ics.uci.edu

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

1



It's time again for....

Friday Nights with Databases



Brought to you by...

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

2

Announcements



❖ Homework notes

- HW #3 was due today (tomorrow if one day late)
- HW #4 will come out on Monday (after the exam)

❖ Exam notes (time flies!)

- **Midterm #1 is on Monday (in class)!**
- We'll use assigned seating – come early!
- Bring an 8.5" x 11" (2-sided) cheat sheet if you want!
- Don't bank on last-minute Piazza answers...

❖ Today's lecture plan:

- Relational languages continued...
- **Note:** Today's info won't be on Monday's exam! ☺

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

3

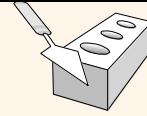
Ex: Wisconsin Sailing Club Database

Sailors				Reserves			Boats		
sid	sname	rating	age	sid	bid	date	bid	bname	color
22	Dustin	7	45.0	22	101	10/10/98	101	Interlake	blue
29	Brutus	1	33.0	22	102	10/10/98	102	Interlake	red
31	Lubber	8	55.5	22	103	10/8/98	103	Clipper	green
32	Andy	8	25.5	22	104	10/7/98	104	Marine	red
58	Rusty	10	35.0	31	102	11/10/98			
64	Horatio	7	35.0	31	103	11/6/98			
71	Zorba	10	16.0	31	104	11/12/98			
74	Horatio	9	35.0	64	101	9/5/98			
85	Art	4	25.5	64	102	9/8/98			
95	Bob	3	63.5	74	103	9/8/93			

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

4

REVIEW: Find names of sailors who've reserved a red boat



Sailors(sid, sname, rating, age) Reserves(sid, bid, day)
 Boats(bid, bname, color)

- ❖ Information about boat color only available in Boats; so need to do two joins:

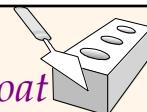
$$\pi_{sname}((\sigma_{color='red'} \text{Boats}) \bowtie \text{Reserves} \bowtie \text{Sailors})$$

- ❖ A more “efficient” solution:

$$\pi_{sname}(\pi_{sid}((\pi_{bid} \sigma_{color='red'} \text{Boats}) \bowtie \text{Res}) \bowtie \text{Sailors})$$

A query optimizer will find the latter, given the 1st query!

Find sailors who've reserved a red or a green boat



Sailors(sid, sname, rating, age) Reserves(sid, bid, day)
 Boats(bid, bname, color)

- ❖ Can identify all red or green boats, then find sailors who've reserved one of these boats:

$$\rho(Tmpboats, (\sigma_{color='red' \vee color='green'} \text{Boats}))$$

$$\pi_{sname}(Tmpboats \bowtie \text{Reserves} \bowtie \text{Sailors})$$

- ❖ Can also define Tempboats using union! (Q: How?)
- ❖ What happens if \vee is replaced by \wedge in this query?

or直接换成and的话，是指这个color是红又是绿，没有这种颜色，因为是想说这个sailor有红色的船又有绿色的船

Find sailors who've reserved a red and a green boat

Sailors(sid, sname, rating, age) Reserves(sid, bid, day)
 Boats(bid, bname, color)

- ❖ Previous approach won't work! Must identify sailors who've reserved red boats and sailors who've reserved green boats, then find their intersection (notice that *sid* is a key for Sailors!):

$$\rho (Tempred, \pi_{sid} ((\sigma_{color='red'} Boats) \bowtie Reserves))$$

$$\rho (Tempgreen, \pi_{sid} ((\sigma_{color='green'} Boats) \bowtie Reserves))$$

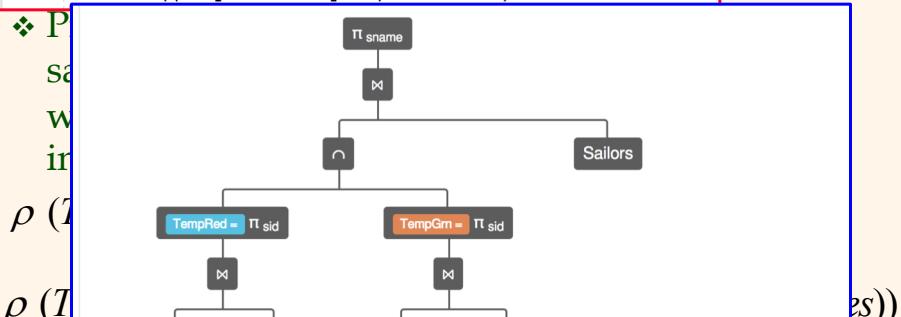
$$\pi_{sname} ((Tempred \cap Tempgreen) \bowtie Sailors)$$

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

7

Find sailors who've reserved a red and a green boat

```
1 TempRed = π sid (σ color = 'red' (Boats) ⋈ Reserves)
2 TempGrn = π sid (σ color = 'green' (Boats) ⋈ Reserves)
3 π sname (((TempRed ∩ TempGrn) ⋈ Sailors))
```



Database

$\pi_{sname} (((\pi_{sid} (\sigma_{color='red'} \text{Boats}) \bowtie \text{Reserves})) \cap ((\pi_{sid} (\sigma_{color='green'} \text{Boats}) \bowtie \text{Reserves}))) \bowtie \text{Sailors})$

8

Find the names of sailors who've reserved all boats

Sailors(sid, sname, rating, age) Reserves(sid, bid, day)
 Boats(bid, bname, color)

- ❖ Uses **division**; schemas of the input relations feeding the / operator must be *carefully chosen*:

$$\rho (Tempsids, (\pi_{sid.bid} Reserves) / (\pi_{bid} Boats))$$

$\pi_{sname} (Tempsids \bowtie Sailors)$ 这两个应该 be the same type

- ❖ To find sailors who've reserved all ‘Interlake’ boats:

$$\dots / \pi_{bid} (\sigma_{bname = 'Interlake'} Boats)$$

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

9

Find the names of sailors who've reserved all boats

RelaX - relational algebra calculator 0.18.2

- ❖ U
fe
 - ❖ To
- | Wisconsin Sailing C... | |
|------------------------|-----------------------------------------------------------|
| Sailors | sid number
sname string
rating number
age number |
| Boats | bid number
bname string
color string |
| Reserves | sid number
bid number
date date |

Relational Algebra SQL

 $\pi \sigma \rho \leftarrow \tau \gamma \wedge \vee \neg = \neq \geq \leq \cap \cup$

```

1 SBpairs =  $\pi_{sid, bid} (Reserves)$ 
2 B =  $\pi_{bid} (Boats)$ 
3 AllSlist = SBpairs  $\div B$ 
4  $\pi_{sname} (AllSlist \bowtie Sailors)$ 

```

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

10

Find the names of sailors who've reserved all boats

RelaX - relational algebra

```
❖ Use fe
❖ To bid
```

Sailors

```
sid number
    fname string
    rating number
    age number
```

Boats

```
bid number
    bname string
    color string
```

Reserves

```
sid number
    bid number
    date date
```

► execute query

```


$$\pi_{\text{fname}} (\text{Sailors} \bowtie (\text{AllList} \div (\text{Reserves} \bowtie \text{Boats})))$$


$$\text{AllList} = \pi_{\text{sid}, \text{bid}}(\text{Reserves})$$


$$B = \pi_{\text{bid}}(\text{Boats})$$


```

$\pi_{\text{fname}} ((\pi_{\text{sid}, \text{bid}}(\text{Reserves})) \div (\pi_{\text{bid}}(\text{Boats}))) \bowtie \text{Sailors}$

Sailors.fname
Dustin

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

11

PS: RelaX Renaming Example...

Relational Algebra

```

 $\rho \sigma \rho \leftarrow \tau \gamma \wedge$ 
 $\rho \text{ sailor\_id} \leftarrow$ 

```

► execute selection

```


$$\rho_{\text{sailor\_id} \leftarrow \text{sid}, \text{sailor\_name} \leftarrow \text{fname}} (\text{Sailors} \bowtie \text{Reserves})$$


```

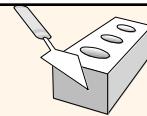
$\rho_{\text{sailor_id} \leftarrow \text{sid}, \text{sailor_name} \leftarrow \text{fname}} (\text{Sailors} \bowtie \text{Reserves})$

Sailors.sailor_id	Sailors.sailor_name	Sailors.rating	Sailors.age	Reserves.bid	Reserves.date
22	Dustin	7	45	101	1998-10-10
22	Dustin	7	45	102	1998-10-10
22	Dustin	7	45	103	1998-10-08
22	Dustin	7	45	104	1998-10-07
31	Lubber	8	55.5	102	1998-10-11
31	Lubber	8	55.5	103	1998-11-06
31	Lubber	8	55.5	104	1998-11-12
64	Horatio	7	35	101	1998-09-05
64	Horatio	7	35	102	1998-09-02
74	Horatio	9	35	103	1993-09-08

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

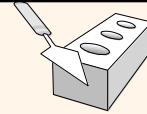
12

Relational Algebra Summary



- ❖ The relational model has (several) rigorously defined query languages that are both simple and powerful in nature.
- ❖ Relational algebra is more “operational”; very useful as an internal representation for query evaluation plans. (SQL “EXPLAIN”)
- ❖ Several ways of expressing a given query; a query optimizer should choose the most efficient version. (Take CS122C...! ☺)
- ❖ We’ll add a few more operators later on...
- ❖ Next up for now: *Relational Calculus*

NEXT: Relational Calculus!



- ❖ Comes in two flavors: *Tuple relational calculus* (TRC) and *Domain relational calculus* (DRC).
- ❖ Calculus has *variables*, *constants*, *comparison ops*, *logical connectives* and *quantifiers*.
 - *TRC*: Variables range over (i.e., get bound to) *tuples*.
 - *DRC*: Variables range over *domain elements* (= field values).
 - Both TRC and DRC are simple subsets of first-order logic.
- ❖ Expressions in the calculus are called *formulas*. An answer tuple is essentially an assignment of constants to variables that make the formula evaluate to *true*.
- ❖ TRC is the basis for various query languages (Quel, SQL, OQL, XQuery, ...), while DRC is the basis for example-based relational query UIs. We’ll study *TRC*!

2. 查询理论

2.1 SQL之母：Relational Calculus

其实SQL并非起源与我们熟知的关系代数（Relational Algebra），实际上关系代数更多地应用于查询的分析和执行。SQL真正之母是很少听说过的关系演算（Relational Calculus），它的查询形式是基于谓词逻辑（Predicate logic）来定义的。具体来说，有元组关系演算（TRC）和域关系演算（DRC）两种：

2.2 Tuple Relational Calculus (TRC) 元组关系演算

TRC被Codd引入作为关系模型的一部分，用来为关系模型提供一种声明式的数据库查询语言。TRC查询的形式为： $\{T \mid \text{Condition}(T)\}$ 。其中T叫做目标（Target），是遍历所有元组的变量。Condition是查询的条件，用具体元组替换T来判断真值。例如， $\{S \mid \text{Staff}(S) \wedge S.\text{salary} > 10000\}$ 。

TRC与SQL有什么对应关系吗？例如上面的例子，S可以看做是SQL中的SELECT部分，但TRC只能返回整个Tuple。Staff(S)可以看做是FROM，而 $S.\text{salary} > 10000$ 可以看做是WHERE。执行方式为：用具体的Tuple去替换S（不能凭空用任意值去替换S吧？没有SQL易理解），看是否满足S在Staff中和 $S.\text{salary}$ 大于10000两个条件。

来看一个复杂一些的TRC，这里就要用到前面逻辑中的概念了。下面S是Free变量，T被量词限定所以是Bound变量。在TRC中，Bound变量用在条件中，用来判断关于Tuple的真假，决定是否保留。而Free变量则用在Target中，用来指定要返回的Tuple。条件中只能有一个Free变量，否则就没法判断真假了。

$$\{S|Student(S) \text{ AND } (\exists T \in Transcript \\ (S.Id = T.StudId \text{ AND } T.CrsCode = 'AAA'))\}$$

转换很简单，等价的SQL就是：

```
SELECTS  
FROM StudentS, TranscriptT  
WHERE S.Id = T.StudId AND T.CrsCode = 'AAA'
```

上面这样一对比就能看出，本质上，**SQL不过是TRC的一层语法糖**：

- **量词哪去了？**：SQL制定了一些约定（Convention）简化了写法，例如SELECT中没有出现的关系都隐式地加上 \exists 量词限定
- **运算丰富**：SQL还融入了聚合函数、关系代数以及集合的交并集运算等（只不过是更多的Sugar）
- **局限性**：但SQL中不允许使用 \forall 量词及否定（Negation）。

所以说，TRC更加通用，但SQL更加简单。

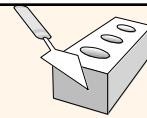
2.3 Domain Relational Calculus (DRC)

DRC的Target部分是由域变量 (Domain Variable) 而非TRC中的整个Tuple组成，即Tuple某个属性的值域中对应的变量。DRC查询的形式为： $\{X_1, X_2, \dots, X_n | condition(X_1, X_2, \dots, X_n)\}$ 。其中每个X都要么是域变量要么是常量。例如，TRC查询 $\{S | Staff(S) \wedge S.\text{salary} > 10000\}$ ，在DRC中可以写作 $\{S.\text{id} | Staff(S.\text{id}, S.\text{salary}) \wedge S.\text{salary} > 10000\}$

因为以属性而非整个元组为Target，所以查询条件中经常出现非Target属性的大量 \exists 也是DRC的标志，但DRC能自动匹配属性值和同名属性关联。所以，某些查询用TRC表达简单，而有些用DRC能更简单些。

TRC和DRC可以相互转换，但转换成关系代数却不容易，有些查询甚至无法转换。

Tuple Relational Calculus



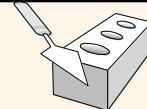
❖ *Query* in TRC has the form:

$$\{ t(\text{attrlist}) \mid P(t) \}$$

❖ *Answer* includes all tuples t with (optionally) specified schema (attrlist) that cause formula $P(t)$ to be *true*.

❖ *Formula* is recursively defined, starting with simple *atomic formulas* (getting tuples from relations or making comparisons of values), and building up bigger and better Boolean formulas using *logical connectives*.

TRC Formulas



❖ *Atomic formula*:

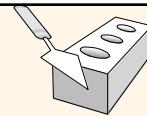
- $r \in R$, or $r \notin R$, or $r.a \text{ op } s.b$, or $r.a \text{ op } \text{constant}$
- op is one of $<$, $>$, \leq , \geq , \neq , $=$

❖ *Formula*:

- an atomic formula, or
- $\neg P$, $P \wedge Q$, $P \vee Q$, where P and Q are formulas, or
- $\exists r \in R (P(r))$, where variable r is *free* in $P(\dots)$, or
- $\forall r \in R (P(r))$, where variable r is *free* in $P(\dots)$, or
- $P \Rightarrow Q$ (pronounced “implies”, equivalent to $(\neg P) \vee Q$)

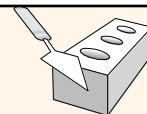
$\exists \nexists \forall \in \notin \neg \wedge \vee \Rightarrow = \neq < > \leq \geq$

Free and Bound Variables



- ❖ The use of a **quantifier** such as $\exists t \in T$ or $\forall t \in T$ in a formula is said to **bind** t .
 - A variable that is **not bound** is **free**.
- ❖ Now let us revisit the definition of a **TRC query**:
 - $\{ t(a_1, a_2, \dots) \mid P(t) \}$
- ❖ There is an important restriction: the variable t that appears to the left of the $|$ ("such that") symbol must be the **only free variable** in the formula $P(\dots)$.
- ❖ Let's look at some examples!

Find sailors with a rating above 7



Sailors(sid, sname, rating, age) Reserves(sid, bid, day)
Boats(bid, bname, color)

$\{ s \mid s \in \text{Sailors} \wedge s.\text{rating} > 7 \}$

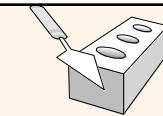
- ❖ This is equivalent to the more general form:

$\{ t(\text{id}, \text{nm}, \text{rtg}, \text{age}) \mid \exists s \in \text{Sailors}$
 $(t.\text{id} = s.\text{sid} \wedge t.\text{nm} = s.\text{sname}$
 $\wedge t.\text{rtg} = s.\text{rating} \wedge t.\text{age} = s.\text{age}$
 $\wedge s.\text{rating} > 7) \}$

(Q: See how each one specifies the answer's schema and values...?)

Note that the second one's schema is different, as we've specified it.)

Find ids of sailors who are older than 30.0 or who have a rating under 8 and are named "Horatio"



Sailors(sid, sname, rating, age) Reserves(sid, bid, day)
Boats(bid, bname, color)

$$\{ t(sid) \mid \exists s \in \text{Sailors} ((s.age > 30.0 \vee (s.rating < 8 \wedge s.sname = \text{"Horatio"}) \wedge t.sid = s.sid) \}$$

❖ Things to notice:

- Again, how result schema and values are specified
- Use of Boolean formula to specify the query constraints
- *Highly declarative nature of this form of query language!*

Ex: TRC Query Semantics

Sailors

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	4	25.5
95	Bob	3	63.5

Reserves

Boats

Sailors

Reserves

Boats

To Be Continued...

