




Introduction to Data Management



Lecture #14 (Relational Languages IV)

Instructor: Mike Carey
mjcarey@ics.uci.edu

It's time again for....

*Friday Nights
With Databases*

Brought to you by...



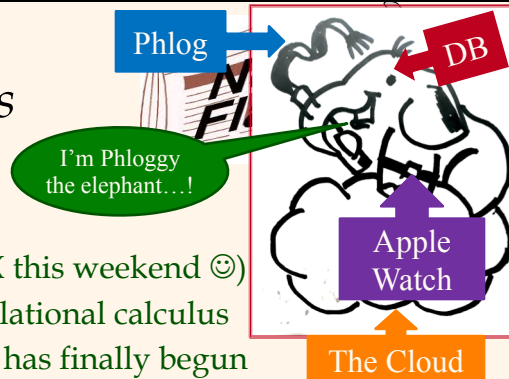
Announcements

❖ HW#4 is in flight...!

- Due Monday (RelaX this weekend ☺)
- Finished with the relational calculus
- Our SQL adventure has finally begun

❖ HW#5 due out Monday ("Monday mode" for now)

- First of a series of SQL-based HW assignments
- Critical that you resolve any MySQL issues! (Take your machine to discussion section, post questions on Piazza, whatever it takes – else you won't survive...!)
- *We now have a project logo...* ☺



Example Data in MySQL

Sailors

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	4	25.5
95	Bob	3	63.5
101	Joan	3	NULL
107	Johan...	NULL	35.0


Reserves

sid	bid	date
22	101	1998-10-10
22	102	1998-10-10
22	103	1998-10-08
22	104	1998-10-07
31	102	1998-11-10
31	103	1998-11-06
31	104	1998-11-12
64	101	1998-09-05
64	102	1998-09-08
74	103	1998-09-08
NULL	103	1998-09-09
1	NULL	2001-01-11
1	NULL	2002-02-02

Boats

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

Find sid's of sailors who've reserved a red or a green boat



Sailors(sid, sname, rating, age)
Reserves(sid, bid, day)
Boats(bid, bname, color)

- ❖ If we replace **OR** by **AND** in this first version, what do we get?
- ❖ **UNION**: Can be used to compute the union of any two *union-compatible* sets of tuples (which are themselves the result of SQL queries).
- ❖ Also available: **EXCEPT** (What would we get if we replaced **UNION** by **EXCEPT**?)

```
SELECT DISTINCT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
AND (B.color='red' OR B.color='green')
```


```
(SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
AND B.color='red')
```

```
UNION
(SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
AND B.color='green')
```

[Note: MySQL vs. Relax - and why?]

SQL vs. TRC

Find sid's of sailors who've reserved a red or a green boat



Sailors(sid, sname, rating, age)
Reserves(sid, bid, day)
Boats(bid, bname, color)

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
AND (B.color='red' OR B.color='green')
```

$\{ t(sid) \mid \exists s \in \text{Sailors} (t.sid = s.sid \wedge$ 这个需要一直括号里面这样去链接两个table
 $\exists r \in \text{Reserves} (r.sid = s.sid \wedge$
 $\exists b \in \text{Boats} (b.bid = r.bid \wedge$
 $(b.color = 'red' \vee b.color = 'green')))) \}$

Find sid's of sailors who've reserved a red and a green boat

Sailors(sid, sname, rating, age)

Reserves(sid, bid, day)

Boats(bid, bname, color)

- ❖ **INTERSECT**: Can be used to compute the intersection of any two *union-compatible* sets of tuples.

- ❖ Included in the SQL/92 standard, but **not** in all systems (incl. MySQL).

- ❖ Contrast symmetry of the UNION and INTERSECT queries with how much the other versions differ.

```
SELECT S.sid
FROM Sailors S, Boats B1, Reserves R1,
     Boats B2, Reserves R2
WHERE S.sid=R1.sid AND R1.bid=B1.bid
     AND S.sid=R2.sid AND R2.bid=B2.bid
     AND (B1.color='red' AND B2.color='green')
```

这里用了两次一样的两个东西，所以可以认为是一个在找红色的，另一个在找绿色的

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
     AND B.color='red'
```

Key field!

```
INTERSECT
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid
     AND B.color='green'
```

Nested Queries


Find names of sailors who've reserved boat #103:

```
SELECT S.sname
FROM Sailors S
WHERE S.sid IN (SELECT R.sid
               FROM Reserves R
               WHERE R.bid=103)
```

WHERE, from, having 都可以用in

- ❖ A very powerful feature of SQL: a WHERE clause can itself contain an SQL query! (Actually, so can SQL's FROM and HAVING clauses!)
- ❖ To find sailors who've *not* reserved #103, use NOT IN.
- ❖ To understand semantics (including cardinality) of nested queries, think nested loops evaluation: *For each Sailors tuple, check qualification by computing subquery.*

Nested Queries *with Correlation*



Find names of sailors who've reserved boat #103:


```
SELECT S.sname
FROM Sailors S
WHERE EXISTS (SELECT *
              FROM Reserves R
              WHERE R.bid=103 AND S.sid=R.sid)
```



- ❖ **EXISTS** is another set comparison operator, like **IN**.
- ❖ Illustrates why, in general, subquery must be re-computed for each Sailors tuple (conceptually).

NOTE: Recall that there was a join way to express this query, too. Relational query optimizers will try to unnest queries into joins when possible to avoid nested loop query evaluation plans.

More on Set-Comparison Operators




- ❖ We've already seen **IN** and **EXISTS**.. Can also use **NOT IN** and **NOT EXISTS**.
- ❖ Also available: *op ANY, op ALL* (for ops: <, >, ≤, ≥, =, ≠)
- ❖ Find sailors whose rating is greater than that of some sailor called Horatio:

```
SELECT *
FROM Sailors S
WHERE S.rating > ANY (SELECT S2.rating
                     FROM Sailors S2
                     WHERE S2.sname='Horatio')
```

So let's try ...
... running w/**ANY** on MySQL
... running w/**ALL** on MySQL

Rewriting INTERSECT Queries Using IN




Find sid's of sailors who've reserved both a red and a green boat:

```
SELECT S.sid
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid AND B.color='red'
      AND S.sid IN (SELECT S2.sid
                    FROM Sailors S2, Boats B2, Reserves R2
                    WHERE S2.sid=R2.sid AND R2.bid=B2.bid
                      AND B2.color='green')
```

- ❖ Similarly, EXCEPT queries can be re-written using NOT IN.
- ❖ This is what you'll need to do when using MySQL (but you can play with **RelaX** for the other set ops).

Division, SQL Style



Find sailors who've reserved all boats.

```
(1) SELECT S.sname
     FROM Sailors S
     WHERE NOT EXISTS
           ((SELECT B.bid
            FROM Boats B)
          EXCEPT
           (SELECT R.bid
            FROM Reserves R
            WHERE R.sid=S.sid))
```

(This Sailor's unreserved Boat ids...!)

Sailors S such that ...

the set of all Boat ids ...

minus ...

this Sailor's reserved Boat ids...

is empty!

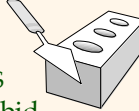
Division in SQL (cont.)

(1)

Find sailors who've reserved all boats.

- ❖ Let's do it the hard(er) way, i.e., without EXCEPT:

```
SELECT S.sname
FROM Sailors S
WHERE NOT EXISTS
  ((SELECT B.bid
    FROM Boats B)
  EXCEPT
  (SELECT R.bid
    FROM Reserves R
    WHERE R.sid=S.sid))
```



(2) SELECT S.sname

FROM Sailors S

WHERE NOT EXISTS (SELECT B.bid
FROM Boats B

Sailors S such that ...

there is no boat B without ...

a Reserves tuple saying that S reserved B

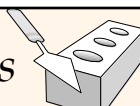
WHERE NOT EXISTS (SELECT R.bid
FROM Reserves R
WHERE R.bid=B.bid
AND R.sid=S.sid))

*This way is **not** that **non-easy**
to understand – right...? (☺)*

Ordering and/or Limiting Query Results

Find the ratings, ids, names, and ages of the three best sailors

```
SELECT S.rating, S.sid, S.sname, S.age
FROM Sailors S
ORDER BY S.rating DESC
LIMIT 3
```

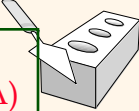


- ❖ The general syntax for this:

```
SELECT [DISTINCT] expressions
FROM tables
[WHERE condition]
....
[ORDER BY expression [ ASC | DESC ]]
LIMIT number_rows [ OFFSET offset_value ];
```

Aggregate Operators

- ❖ Significant extension of the relational algebra.



```
COUNT(*)
COUNT( [DISTINCT] A)
SUM( [DISTINCT] A)
AVG( [DISTINCT] A)
MAX(A)
MIN(A)
```

single column

```
SELECT COUNT(*)  总共有多少条
FROM Sailors S
```

```
SELECT AVG(S.age)
FROM Sailors S
WHERE S.rating=10
```

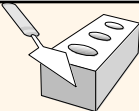
```
SELECT S.sname
FROM Sailors S
WHERE S.rating= (SELECT MAX(S2.rating)
                  FROM Sailors S2)
```

```
SELECT COUNT(DISTINCT S.rating)
FROM Sailors S
WHERE S.sname= 'Bob'
```

```
SELECT AVG(DISTINCT S.age)
FROM Sailors S
WHERE S.rating=10
```

Find name and age of the oldest sailor(s)

- ❖ That first try is *illegal*!
(We'll see why shortly,
when we do **GROUP BY**.)
- ❖ *Nit*: The third version is
equivalent to the second
one, and is allowed in the
SQL/92 standard, but not
supported in all systems.




```
SELECT S.sname, MAX(S.age)
FROM Sailors S
```

```
SELECT S.sname, S.age
FROM Sailors S
WHERE S.age =
      (SELECT MAX(age)
       FROM Sailors)
```

```
SELECT S.sname, S.age
FROM Sailors S
WHERE (SELECT MAX(S2.age)
       FROM Sailors S2)
      = S.age
```


Motivation for Grouping




- ❖ So far, we've applied aggregate operators to all (qualifying) tuples. Sometimes, we want to apply them to each of several *groups* of tuples.
- ❖ Consider: *Find the age of the youngest sailor for each rating level.*
 - In general, we don't know how many rating levels exist, and what the rating values for these levels are!
 - Suppose we know that rating values go from 1 to 10; we can write 10 queries that look like this (☺):

For $i = 1, 2, \dots, 10$:

```
SELECT MIN(S.age)
FROM Sailors S
WHERE S.rating = i
```


Queries With GROUP BY and HAVING



```
SELECT    [DISTINCT] target-list
FROM      relation-list
WHERE     qualification
GROUP BY  grouping-list
HAVING    group-qualification
```


- ❖ The *target-list* contains (i) attribute names and (ii) terms with aggregate operations (e.g., MIN (S.age)).
 - The attribute list (i) must be a subset of *grouping-list*. Intuitively, each answer tuple corresponds to a *group*, and these attributes must have a **single value per group**. (A *group* is a set of tuples that have the same value for all attributes in *grouping-list*.)

Conceptual Evaluation



- ❖ The cross-product of *relation-list* is computed, tuples that fail the *qualification* are discarded, 'unnecessary' fields are deleted, and the remaining tuples are partitioned into groups by the value of attributes in *grouping-list*.
- ❖ A *group-qualification* (HAVING) is then applied to eliminate some groups. Expressions in *group-qualification* must also have a single value per group!
 - In effect, an attribute in *group-qualification* that is not an argument of an aggregate op must appear in *grouping-list*.
(Note: SQL doesn't consider primary key semantics here.)
- ❖ One answer tuple is generated per qualifying group.

Find age of the youngest sailor with age ≥ 18 for each rating with at least 2 such sailors.



每个group里面的age的平均数

```
SELECT S.rating, MIN(S.age)
      AS minage
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING COUNT(*) >= 2
```


Sailors instance:

sid	sname	rating	age
22	dustin	7	45.0
29	brutus	1	33.0
31	lubber	8	55.5
32	andy	8	25.5
58	rusty	10	35.0
64	horatio	7	35.0
71	zorba	10	16.0
74	horatio	9	35.0
85	art	3	25.5
95	bob	3	63.5
96	frodo	3	25.5

Answer relation:

rating	minage
3	25.5
7	35.0
8	25.5

*Find age of the youngest sailor with age ≥ 18
for each rating with at least 2 such sailors.*



rating	age
7	45.0
1	33.0
8	55.5
8	25.5
10	35.0
7	35.0
10	16.0
9	35.0
3	25.5
3	63.5
3	25.5



rating	age
1	33.0
3	25.5
3	63.5
3	25.5
7	45.0
7	35.0
8	55.5
8	25.5
9	35.0
10	35.0



rating	minage
3	25.5
7	35.0
8	25.5