




# *Introduction to Data Management*



## *Lecture #15 (SQL, the Sequel...)*

Instructor: Mike Carey  
mjcarey@ics.uci.edu

## *Announcements*




### ❖ HW stuff

- Hopefully everyone's feeling Relaxed today. ☺
- HW4's solution will appear tomorrow at 5pm.
- There will be no actual TRC HW (just quiz stuff).

### ❖ Other logistical stuff

- Midterm grading is still in progress...
- We will now continue our SQL adventure!

## Grouped Aggregation Queries (Review)



```
SELECT    [DISTINCT] target-list
FROM      relation-list
WHERE     qualification
GROUP BY  grouping-list
HAVING    group-qualification
```

For example...


```
SELECT S.rating, MIN(S.age) AS minage
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING COUNT(*) >= 2
```

Group aggregate(s)

Grouping field(s)

Group predicate(s)

Find age of the youngest sailor with age  $\geq 18$  for each rating with at least 2 such sailors. (Review)



Query:

```
SELECT S.rating, MIN(S.age) AS minage
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING COUNT(*) >= 2
```


Sailors instance:

sid	sname	rating	age
22	dustin	7	45.0
29	brutus	1	33.0
31	lubber	8	55.5
32	andy	8	25.5
58	rusty	10	35.0
64	horatio	7	35.0
71	zorba	10	16.0
74	horatio	9	35.0
85	art	3	25.5
95	bob	3	63.5
96	frodo	3	25.5

Answer relation:

rating	minage
3	25.5
7	35.0
8	25.5

## Example Data in MySQL



**Sailors**

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	4	25.5
95	Bob	3	63.5
101	Joan	3	NULL
107	Johan...	NULL	35.0


**Reserves**

sid	bid	date
22	101	1998-10-10
22	102	1998-10-10
22	103	1998-10-08
22	104	1998-10-07
31	102	1998-11-10
31	103	1998-11-06
31	104	1998-11-12
64	101	1998-09-05
64	102	1998-09-08
74	103	1998-09-08
NULL	103	1998-09-09
1	NULL	2001-01-11
1	NULL	2002-02-02

**Boats**

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

*For each red boat, find the number of reservations for this boat*




for each的就是用group

```
SELECT B.bid, COUNT(*) AS scount
FROM Sailors S, Boats B, Reserves R
WHERE S.sid=R.sid AND R.bid=B.bid AND B.color='red'
GROUP BY B.bid
```

- ❖ We're grouping over a join of three relations!
- ❖ What do we get if we remove *B.color= 'red'* from the WHERE clause and add a HAVING clause with this condition? (Hint: Trick question... ☺)
- ❖ What if we drop Sailors and the condition involving S.sid?


*Find age of the youngest sailor with age > 18  
for each rating with at least 2 sailors (of **any** age)*



```
SELECT S.rating, MIN(S.age)
FROM Sailors S
WHERE S.age > 18
GROUP BY S.rating
HAVING 1 < (SELECT COUNT(*)
            FROM Sailors S2
            WHERE S.rating=S2.rating)
```

- ❖ Shows HAVING clause can also contain a subquery.
- ❖ Compare this with the query where we considered only ratings with 2 or more sailors over 18!
- ❖ What if HAVING clause were replaced by:
  - HAVING COUNT(\*) >1

*Find those ratings and average ages for which  
the average Sailor age is the minimum age over  
all of the Sailors*



- ❖ Aggregate operations can't be nested! (**WRONG...**)

```
SELECT S.rating
FROM Sailors S
WHERE S.age = (SELECT MIN(AVG(S2.age)) FROM Sailors S2)
```


- ❖ A correct solution (in SQL/92):

```
SELECT Temp.rating, Temp.avgage
FROM (SELECT S.rating, AVG(S.age) AS avgage
      FROM Sailors S
      GROUP BY S.rating) AS Temp
WHERE Temp.avgage = (SELECT MIN(age) FROM Sailors)
```

Compute the  
average age for  
each rating...

Find the overall  
minimum age

## SQL's WITH Clause



*Ex: Find those ratings and average ages for which the average Sailor age is the minimum age over all of the Sailors*


❖ Our first solution was:

```
SELECT Temp.rating, Temp.avgage
FROM (SELECT S.rating, AVG(S.age) AS avgage
      FROM Sailors S
      GROUP BY S.rating) AS Temp
WHERE Temp.avgage = (SELECT MIN(age) FROM Sailors)
```

❖ We could use a **WITH** clause here for clarity:

```
WITH Temp AS (SELECT S.rating, AVG(S.age) AS avgage
              FROM Sailors S
              GROUP BY S.rating)
SELECT Temp.rating, Temp.avgage
FROM Temp
WHERE Temp.avgage = (SELECT MIN(age) FROM Sailors)
```

## Null Values in SQL



❖ Field values in a tuple are sometimes *unknown* (e.g., a rating has not yet been assigned) or *inapplicable* (e.g., there is no spouse's name).

- SQL provides the special value null for such situations.

❖ The presence of *null* complicates many issues. E.g.:

- Special operators needed to check if value is/is not *null*.
- Is *rating* > 8 true or false when *rating* is equal to *null*? What about **AND**, **OR** and **NOT** connectives?
- We need a 3-valued logic (true, false and *unknown*).
- Meaning of constructs must be defined carefully. (The **WHERE** clause eliminates rows that don't evaluate to *true*.)
- New operators (in particular, *outer joins*) possible/needed.

## Ex: Sailors w/Some Null Values

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	4	25.5
95	Bob	3	63.5
101	Joan	3	NULL
107	Johannes	NULL	35.0

*Q:* Which kind(s) of *null* are each of these null values?

## Nulls and SQL's 3-Valued Logic

AND	true	false	unknown
true	true	false	unknown
false	false	false	false
unknown	unknown	false	unknown


OR	true	false	unknown
true	true	true	true
false	true	false	unknown
unknown	true	unknown	unknown

### NOT

true	false
false	true
unknown	unknown

*Note:* SQL arithmetic expressions involving *null* values will yield *null* values (Ex: EMP.sal + EMP.bonus)

## Basic SQL Queries w/Nulls




sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	4	25.5
95	Bob	3	63.5
101	Joan	3	NULL
107	Johannes	NULL	35.0

```
SELECT *
FROM Sailors S
WHERE age > 35.0
```

```
SELECT *
FROM Sailors S
WHERE age <= 35.0
```

```
SELECT COUNT(*)
FROM Sailors S
WHERE age > 35.0
      OR age <= 35.0
      OR age IS NULL
```


## Ex: Sailors and Reserves w/Nulls



sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	4	25.5
95	Bob	3	63.5
101	Joan	3	NULL
107	Johannes	NULL	35.0

sid	bid	date
22	101	1998-10-10
22	102	1998-10-10
22	103	1998-10-08
22	104	1998-10-07
31	102	1998-11-10
31	103	1998-11-06
31	104	1998-11-12
64	101	1998-09-05
64	102	1998-09-08
74	103	1998-09-08
NULL	103	1998-09-09
1	NULL	2001-01-11
1	NULL	2002-02-02

## Nulls w/Aggregates



sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	4	25.5
95	Bob	3	63.5
101	Joan	3	NULL
107	Johannes	NULL	35.0


SELECT COUNT(rating)  
FROM Sailors (11) Null不算

SELECT  
COUNT (DISTINCT rating)  
FROM Sailors (7)

SELECT SUM(rating),  
COUNT(rating),  
AVG(rating)  
FROM Sailors  
(70, 11, 6.3636) Null不算

(Useful, but logically "wrong"!)

## Nulls w/Aggregates & Grouping



sid	bid	date
22	101	1998-10-10
22	102	1998-10-10
22	103	1998-10-08
22	104	1998-10-07
31	102	1998-11-10
31	103	1998-11-06
31	104	1998-11-12
64	101	1998-09-05
64	102	1998-09-08
74	103	1998-09-08
NULL	103	1998-09-09
1	NULL	2001-01-11
1	NULL	2002-02-02

SELECT COUNT( DISTINCT bid)  
FROM Reserves (4) grouping会有

SELECT bid, COUNT(\*)  
FROM Reserves  
GROUP BY bid



bid	COUNT(*)
NULL	2
101	2
102	3
103	4
104	2



Some “dangling” tuple examples


## Nulls w/Joins → Inner vs. Outer Joins

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	4	25.5
95	Bob	3	63.5
101	Joan	3	NULL
107	Johannes	NULL	35.0

sid	bid	date
22	101	1998-10-10
22	102	1998-10-10
22	103	1998-10-08
22	104	1998-10-07
31	102	1998-11-10
31	103	1998-11-06
31	104	1998-11-12
64	101	1998-09-05
64	102	1998-09-08
74	103	1998-09-08
NULL	103	1998-09-09
1	NULL	2001-01-11
1	NULL	2002-02-02

## Inner vs. Outer Joins in SQL

```
SELECT DISTINCT s.sname, r.date
FROM Sailors s, Reserves r
WHERE s.sid = r.sid
```



sname	date
Dustin	1998-10-10
Dustin	1998-10-08
Dustin	1998-10-07
Lubber	1998-11-10
Lubber	1998-11-06
Lubber	1998-11-12
Horatio	1998-09-05
Horatio	1998-09-08

## SQL中INNER JOIN、LEFT JOIN、RIGHT JOIN、FULL JOIN区别

sql中的连接查询有inner join(内连接)、left join(左连接)、right join(右连接)、full join(全连接)四种方式，它们之间其实并没有太大区别，仅仅是查询出来的结果有所不同。例如我们有两张表：

"Persons" 表：

Id_P	LastName	FirstName	Address	City
1	Adams	John	Oxford Street	London
2	Bush	George	Fifth Avenue	New York
3	Carter	Thomas	Changan Street	Beijing

"Orders" 表：

Id_O	OrderNo	Id_P
1	77895	3
2	44678	3
3	22456	1
4	24562	1
5	34764	65

Orders表通过外键Id\_P和Persons表进行关联。

**1.inner join，在两张表进行连接查询时，只保留两张表中完全匹配的结果集。**

我们使用inner join对两张表进行连接查询，sql如下：

```
1 SELECT Persons.LastName, Persons.FirstName, Orders.OrderNo
2 FROM Persons
3 INNER JOIN Orders
4 ON Persons.Id_P=Orders.Id_P
5 ORDER BY Persons.LastName
```

查询结果集：

LastName	FirstName	OrderNo
Adams	John	22456
Adams	John	24562
Carter	Thomas	77895
Carter	Thomas	44678

此种连接方式Orders表中Id\_P字段在Persons表中找不到匹配的，则不会列出来。

**2.left join**,在两张表进行连接查询时，会返回左表所有的行，即使在右表中没有匹配的记录。

我们使用left join对两张表进行连接查询，sql如下：

```
1 SELECT Persons.LastName, Persons.FirstName, Orders.OrderNo
2 FROM Persons
3 LEFT JOIN Orders
4 ON Persons.Id_P=Orders.Id_P
5 ORDER BY Persons.LastName
```

查询结果如下：

LastName	FirstName	OrderNo
Adams	John	22456
Adams	John	24562
Carter	Thomas	77895
Carter	Thomas	44678
Bush	George	

可以看到，左表（Persons表）中LastName为Bush的行的Id\_P字段在右表（Orders表）中没有匹配，但查询结果仍然保留该行。

**3.right join**,在两张表进行连接查询时，会返回右表所有的行，即使在左表中没有匹配的记录。

我们使用right join对两张表进行连接查询，sql如下：

```
1 SELECT Persons.LastName, Persons.FirstName, Orders.OrderNo
2 FROM Persons
3 RIGHT JOIN Orders
4 ON Persons.Id_P=Orders.Id_P
5 ORDER BY Persons.LastName
```

查询结果如下：

LastName	FirstName	OrderNo
Adams	John	22456
Adams	John	24562
Carter	Thomas	77895
Carter	Thomas	44678
		34764

Orders表中最后一条记录Id\_P字段值为65，在左表中没有记录与之匹配，但依然保留。

**4.full join,在两张表进行连接查询时， 返回左表和右表中所有没有匹配的行。**

我们使用full join对两张表进行连接查询， sql如下：

```
1 SELECT Persons.LastName, Persons.FirstName, Orders.OrderNo
2 FROM Persons
3 FULL JOIN Orders
4 ON Persons.Id_P=Orders.Id_P
5 ORDER BY Persons.LastName
```

查询结果如下：

LastName	FirstName	OrderNo
Adams	John	22456
Adams	John	24562
Carter	Thomas	77895
Carter	Thomas	44678
Bush	George	
		34764

查询结果是left join和right join的并集。

这些连接查询的区别也仅此而已。

## Inner vs. Outer Joins in SQL (2)

```
SELECT DISTINCT s.sname, r.date  
FROM Sailors s INNER JOIN Reserves r ON s.sid = r.sid
```

(“INNER” is optional,  
and will be the default  
type of JOIN assumed if  
one isn’t specified)

本来就只显示inner join




sname	date
Dustin	1998-10-10
Dustin	1998-10-08
Dustin	1998-10-07
Lubber	1998-11-10
Lubber	1998-11-06
Lubber	1998-11-12
Horatio	1998-09-05
Horatio	1998-09-08

## Inner vs. Outer Joins in SQL (3)

```
(1) SELECT DISTINCT s.sname, r.date  
FROM Sailors s LEFT OUTER JOIN Reserves r ON s.sid = r.sid  
(2) SELECT DISTINCT s.sname, r.date  
FROM Reserves r RIGHT OUTER JOIN Sailors s ON s.sid = r.sid
```

### ❖ Variations on a theme:

- JOIN (or INNER JOIN)
- LEFT OUTER JOIN
- RIGHT OUTER JOIN
- FULL OUTER JOIN



```
join_table:  
  table_reference [INNER | CROSS] JOIN table_factor [join_condition]  
  | table_reference STRAIGHT_JOIN table_factor  
  | table_reference STRAIGHT_JOIN table_factor ON conditional_expr  
  | table_reference {LEFT|RIGHT} [OUTER] JOIN table_reference join_condition  
  | table_reference NATURAL [INNER | {LEFT|RIGHT} [OUTER]] JOIN table_factor  
  
join_condition:  
  ON conditional_expr  
  | USING (column_list)
```

sname	date
Dustin	1998-10-10
Dustin	1998-10-08
Dustin	1998-10-07
Lubber	1998-11-10
Lubber	1998-11-06
Lubber	1998-11-12
Horatio	1998-09-05
Horatio	1998-09-08
Brutus	NULL
Andy	NULL
Rusty	NULL
Zorba	NULL
Art	NULL
Bob	NULL