




# *Introduction to Data Management*

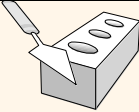


## *Lecture #16 (SQL, the Sequel's Sequel...)*

Instructor: Mike Carey  
mjcarey@ics.uci.edu

# *Announcements*



- 
- ❖ First SQL query HW is now underway
    - Hopefully everyone has MySQL working
    - Get the latest version of the questions! (*Sorry...!* ☹)
  - ❖ Grading is in progress for many things
    - HW #2 should be done any minute (!)
    - Other HW's are in progress in parallel
    - Trying to get Midterm #1 done by week's end
    - (430 is a pretty big number...)

## Example Data in MySQL

**Sailors**

sid	sname	rating	age
22	Dustin	7	45.0
29	Brutus	1	33.0
31	Lubber	8	55.5
32	Andy	8	25.5
58	Rusty	10	35.0
64	Horatio	7	35.0
71	Zorba	10	16.0
74	Horatio	9	35.0
85	Art	4	25.5
95	Bob	3	63.5
101	Joan	3	NULL
107	Johan...	NULL	35.0

**Reserves**

sid	bid	date
22	101	1998-10-10
22	102	1998-10-10
22	103	1998-10-08
22	104	1998-10-07
31	102	1998-11-10
31	103	1998-11-06
31	104	1998-11-12
64	101	1998-09-05
64	102	1998-09-08
74	103	1998-09-08
NULL	103	1998-09-09
1	NULL	2001-01-11
1	NULL	2002-02-02

**Boats**

bid	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

## Inner vs. Outer Joins in SQL (3)

(1) SELECT DISTINCT s.sname, r.date  
FROM Sailors s **LEFT OUTER JOIN** Reserves r **ON** s.sid = r.sid

(2) SELECT DISTINCT s.sname, r.date  
FROM Reserves r **RIGHT OUTER JOIN** Sailors s **ON** s.sid = r.sid

### ❖ Variations on a theme:

- JOIN (or INNER JOIN)
- LEFT OUTER JOIN
- RIGHT OUTER JOIN
- FULL OUTER JOIN


```

join_table:
  table_reference [INNER | CROSS] JOIN table_factor [join_condition]
  | table_reference NATURAL JOIN table_factor
  | table_reference {LEFT|RIGHT} [OUTER] JOIN table_factor join_condition
  | table_reference NATURAL [INNER | {LEFT|RIGHT} [OUTER]] JOIN table_factor

join_condition:
  ON conditional_expr
  | USING (column_list)
  
```

sname	date
Dustin	1998-10-10
Dustin	1998-10-08
Dustin	1998-10-07
Lubber	1998-11-10
Lubber	1998-11-06
Lubber	1998-11-12
Horatio	1998-09-05
Horatio	1998-09-08
Brutus	NULL
Andy	NULL
Rusty	NULL
Zorba	NULL
Art	NULL
Bob	NULL

## An Algebra Side Note...



### ❖ As a side note:

- The underlying operations are also part of the extended relational algebra, which adds...


- Outer joins (left, right, and full)
- Ordering (sorting)
- Grouping (w/aggregates)
- ....

- ❖ You can play around with those extensions on the relational algebra (RelaX) site that you used for the recently completed RA HW (if you're curious)!

- <https://dbis-uibk.github.io/relax/help.htm>

## Updates: Oh **CRUD**\*!

(\***C**reate, **R**etrieve, **U**ppdate, **D**estroy)



- ❖ Can add one or more tuples using INSERT:

```
INSERT INTO Students (sid, name, login, age, gpa)
VALUES (53688, 'Smith', 'smith@ee', 18, 3.2)
```

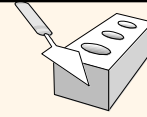
```
INSERT INTO Students (sid, name, login, age, gpa)
SELECT ... (your favorite SQL query goes here!) ...
```

- ❖ Can DELETE all tuples satisfying any SQL query condition:

```
DELETE FROM Students S
WHERE S.sid IN (SELECT X.sid FROM Banned X)
```

## Updates: Oh **CRUD**!

(Cont.)




- ❖ Can change one or more tuples using UPDATE:

```
UPDATE Sailors
SET  sname = 'King Arthur',
     rating = rating + 1
WHERE sname = 'Art';
```

- ❖ A few things to note:

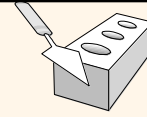
- LHS of SET is column name, RHS is (any) expression
- WHERE predicate is any SQL condition, which again means SQL subqueries are available as a tool, e.g., to search for targets based on multiple tables' content

## SQL Data Integrity (Largely *Review*)



- ❖ An *integrity constraint* describes a condition that every *legal instance* of a relation must satisfy.
  - Inserts/deletes/updates that violate IC's are disallowed.
  - Can be used to ensure application semantics (e.g., *sid* is a key, *bid* refers to a known boat) or prevent inconsistencies (e.g., *sname* has to be a string, integer *age* must be < 120)
- ❖ *Types of IC's*: Domain constraints, primary key constraints, foreign key constraints, unique constraints, general constraints.
  - *Domain constraints*: Field values must be of the right type (i.e., per the schema specification). Always enforced!

## SQL Data Integrity (Cont.)



❖ So far we have been making good use of:

- PRIMARY KEY
- UNIQUE
- NOT NULL
- FOREIGN KEY

*Note:* MySQL with InnoDB actually permits a foreign key to reference any indexed column(s)...

❖ Other features for ensuring field value integrity:

- DEFAULT (alternative to NULL for missing values)
- CHECK (called “general” in the book, kind of...)

❖ More powerful integrity features include

- ~~ASSERTION~~ (called “general” in the book, correctly ☺)
- TRIGGER (a sledge hammer to use when all else fails!)

## Some Integrity Related Examples

CREATE TABLE Sailors

( sid INTEGER,  
sname CHAR(10),  
rating INTEGER,  
age REAL **DEFAULT 18.0**,  
PRIMARY KEY (sid),

**CHECK ( rating >= 1  
AND rating <= 10 ) )**

❖ CHECK is useful when more general ICs than just keys are involved.

❖ Could use SQL subqueries to express richer constraints (if supported ☺).

❖ Constraints can be named (to manage them).

CREATE TABLE Reserves


( sname CHAR(10),  
bid INTEGER,  
day DATE,  
PRIMARY KEY (bid,day),

~~CONSTRAINT noInterlakeRes~~

~~CHECK ( 'Interlake' <>  
(SELECT B.bname  
FROM Boats B  
WHERE B.bid=bid)) )~~


*Note:* Unfortunately, MySQL currently ignores CHECK constraints...

## Enforcing Referential Integrity (RI)



- ❖ Consider Sailors and Reserves; *sid* in Reserves is a foreign key that references Sailors.
- ❖ What should be done if a Reserves tuple with a non-existent sailor id is inserted? (**A: Reject it!**)
- ❖ What should be done if a Sailors tuple is deleted?
  - Also delete all Reserves tuples that refer to it, or
  - Disallow deletion of a Sailors that's being referred to, or
  - Set *sid* in Reserves tuples that refer to it to some *default sid*.
  - (In SQL, could also: Set *sid* in Reserves tuples that refer to it to *null*, denoting 'unknown' or 'inapplicable'.)
- ❖ Similar issue if the primary key of a Sailor is updated.

## RI Enforcement in SQL (Reminder)




- ❖ SQL/92 and SQL:1999 support all 4 options on deletes and updates.
  - Default is **NO ACTION** (delete/update is rejected)
  - **CASCADE** (also delete all tuples that refer to the deleted tuple)
  - **SET NULL / SET DEFAULT** (set foreign key value of referencing tuple)

Ex:

```
CREATE TABLE Reserves
(sid INTEGER,
bid INTEGER,
date DATE,
....
FOREIGN KEY (sid)
REFERENCES Sailors
ON DELETE CASCADE
ON UPDATE SET NULL)
```


*Odd combo; just illustrating what's possible here...*

## Triggers in SQL



- ❖ Trigger: a procedure that runs automatically if specified changes occur to the DBMS
- ❖ Three parts:
  - Event (activates the trigger)
  - Condition (tests if the trigger should run)
  - Action (what happens if the trigger runs)
- ❖ Can be used to do “whatever”!
  - One SQL statement or sequence/flow of statements; can also cause the current update to bail out.
  - Details vary WIDELY from vendor to vendor (!)
  - Major source of “vendor lock-in”, along with the *stored procedure language* (= trigger action language)

## Trigger Syntax (**MySQL**)

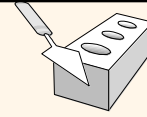


```
CREATE [DEFINER = { user | CURRENT_USER }]
TRIGGER trigger_name
trigger_time trigger_event
ON tbl_name
FOR EACH ROW
[trigger_order]
trigger_body
```

trigger\_time: { BEFORE | AFTER }  
trigger\_event: { INSERT | UPDATE | DELETE }  
trigger\_order: { FOLLOWS | PRECEDES } other\_trigger\_name

<https://dev.mysql.com/doc/refman/8.0/en/trigger-syntax.html>


## Trigger Example (MySQL)



```
DELIMITER $$      -- Necessary to make semicolons great again... ☺  
                  -- (Prevents them from ending the input statement!)
```

```
CREATE TRIGGER youngSailorUpdate  
AFTER INSERT ON Sailors  
FOR EACH ROW      Note: FOR EACH ROW provides less  
                    power than FOR EACH STATEMENT  
                    (e.g., can't compute average new age)  
BEGIN  
    IF NEW.age < 18 THEN  
        INSERT INTO YoungSailors (sid, sname, age, rating)  
        VALUES (NEW.sid, NEW.sname, NEW.age, NEW.rating);  
    END IF;  
END;
```

## Trigger Example (MySQL, cont'd.)




- ☐ INSERT INTO Sailors(sid, sname, rating, age)  
VALUES (777, 'Lucky', 7, 77);
- ☒ INSERT INTO Sailors(sid, sname, rating, age)  
VALUES (778, 'Lucky Jr', 7, 7);

(NOTE: Look at *YoungSailors* table content after each one!)



## Another Trigger Example (*MySQL*)



-- Let's implement a poor man's CHECK constraint!  
DELIMITER \$\$

```
CREATE TRIGGER checkSailorAge
AFTER INSERT ON Sailors
FOR EACH ROW
BEGIN
    IF NEW.age < 18 THEN
        SIGNAL SQLSTATE '02000'
        SET MESSAGE_TEXT =
            'Warning: Sailors can not be under 18!';
    END IF;
END;
```