



Introduction to Data Management

Lecture 20 (Storage and Indexing, cont.)



Instructor: Mike Carey
mjcarey@ics.uci.edu




It's time again for....

Friday Nights with Databases...!

Brought to you by...




Announcements



- ❖ Midterm #2 is **Wednesday (5/22) at 5 PM**
 - Relational languages (see syllabus!)
 - Sample exam from last year is available
 - Assigned seating, similar to last time
- ❖ HW #6 is due on **Monday at 7 PM**
 - One late “day” (*22 hours*) will be available
 - Solution coming Tuesday right after **5 PM** (*really*)
- ❖ Today’s lecture plan
 - More about database indexes
 - (Not on Midterm #2, of course)


Tree-Structured Indexes: Over(*re*)view




- ❖ As for any index, 3 alternatives for data entries (k^*):
 - Record with key k
 - $\langle k, \text{rid of record with key } k \rangle$
 - $\langle k, \text{list of rids of records with key } k \rangle$
- ❖ This data entry choice is orthogonal (\perp) to the *indexing technique* used to locate the data entries.
- ❖ Tree-structured indexing techniques can support both *range searches* and *equality searches*.
- ❖ ISAM: static structure; B+ tree: dynamic, adjusts gracefully under inserts and deletes.

Example ISAM Tree


- ❖ Suppose each node can hold 2 entries (really more like 200, since nodes are disk pages!)



After Inserting 23*, 48*, 41*, 42* ...




... Then Deleting 42*, 51*, 97*



➡ Note that 51* still appears in index levels, but **not** in leaf!


B+ Tree: Most Widely Used Index!

- ❖ Insert/delete at $\log_F N$ cost; keep tree *height-balanced*. (F = fanout, N = # leaf pages)
- ❖ **Minimum 50% occupancy** (except for root).
Each node contains $d \leq m \leq 2d$ entries.
The (mythical) d is called the *order* of the B+ tree.
- ❖ Supports equality and range-searches efficiently.



Example B+ Tree


- ❖ Search begins at root, and key comparisons direct the search to a leaf (as in ISAM).
- ❖ Ex: Search for 5*, 15*, all data entries $\geq 24^*$, ...



➡ Based on the search for 15*, we know it is not in the tree!


B+ Tree: Most Widely Used Index!

- ❖ Insert/delete at $\log_F N$ cost; keep tree *height-balanced*. (F = fanout, N = # leaf pages)
- ❖ Minimum 50% occupancy (except for root). Each node contains $d \leq m \leq 2d$ entries. The (mythical) d is called the *order* of the B+ tree.
- ❖ Supports equality and range-searches efficiently.



Example B+ Tree


- ❖ Search begins at root, and key comparisons direct the search to a leaf (as in ISAM).
- ❖ Ex: Search for 5*, 15*, all data entries $\geq 24^*$, ...



➡ Based on the search for 15*, we know it is not in the tree!

Example B+ Tree (a clarification)

- ❖ Search begins at root, and key comparisons direct the search to a leaf (as in ISAM).
- ❖ Ex: Search for 5*, 15*, all data entries $\geq 24^*$, ...




➡ Based on the search for 15*, we know it is not in the tree!

Inserting a Data Entry into a B+ Tree




- ❖ Find correct leaf L (by *searching* for the new k).
- ❖ Put new data entry (k^* , a.k.a. $(k, I(k))$) in leaf L .
 - If L has enough space, *done!* (Most likely case!)
 - Else, must split L (into L and a new node $L2$)
 - Redistribute entries evenly and copy up middle key.
 - Insert new index entry pointing to $L2$ into parent of L .
- ❖ This can happen recursively.
 - To split an *index node*, redistribute entries evenly but push up the middle key. (Contrast with leaf splits!)
- ❖ Splits “grow” tree; root split increases its height.
 - Tree growth: gets wider or one level taller at top.


Inserting 8^* into Example B+ Tree



- ❖ Observe how minimum occupancy is guaranteed in both leaf and index pg splits.
- ❖ Note difference between *copy-up* and *push-up*; be sure you understand the reasons for this!




Example B+ Tree After Inserting 8*



- ❖ Notice that root was split, leading to increase in height.
- ❖ In this example, could avoid split by redistributing entries; however, not usually done in practice. (Q: Why is that?)

Let's Go Live...! (Demo Time!)



Note (see Piazza): Very cool online B+ tree viz tool available (☺)



- <https://www.cs.usfca.edu/~galles/visualization/BPlusTree.html>
- Only slight differences from our defs (e.g., key 13 above → 14)
- Their “Max. Degree” is our $2d+1$ (limit of 5 ptrs/node above)

Deleting a Data Entry from a B+ Tree



- ❖ Start at root, find leaf L where entry belongs.
- ❖ Remove the entry.
 - If L is still at least half-full, *done!*
 - If L has only $d-1$ entries,
 - Try to **redistribute**, borrowing from sibling (adjacent node with same parent as L).
 - If re-distribution fails, **merge** L and sibling.
- ❖ If merge occurred, must delete search-guiding entry (pointing to L or sibling) from parent of L .
- ❖ Merge could propagate to root, decreasing height.

Example Tree After (Inserting 8^* , Then) Deleting 19^* and 20^* ...



- ❖ Deleting 19^* is easy.
- ❖ Deleting 20^* is done with redistribution.
Notice how middle key is **copied up**.

... And Then Deleting 24*

- ❖ Must merge.
- ❖ Observe “*toss*” of index entry (on right), and “*pull down*” of index entry (below).

