



Introduction to Data Management

Lecture #22 (Physical DB Design)



Instructor: Mike Carey
mjcarey@ics.uci.edu




It's the penultimate episode of....

***Friday Nights with
Databases...!***

Brought to you by...




Announcements




- ❖ No class Monday! (*Awwwww.....*)
- ❖ Two HW assignments remain:
 - HW #7: Due next Thursday, May 30th (5 PM)
 - Physical DB design (for MySQL and beyond)
 - HW #8: Due on Thursday, June 6th (5 PM)!
 - NoSQL (and **NoLateDay** due to Endterm timing)
- ❖ Today's plan :
 - Today: Physical DB design (esp. indexing)
 - Next up: **NoSQL & Big Data** (*a la* AsterixDB)
 - Not in the textbook, so... See the wiki for readings!

Overview




- ❖ After ER design, schema refinement, and the definition of views, we have the *conceptual* and *external* schemas for our database.
- ❖ Next step is to *choose indexes*, make *clustering decisions*, and *refine* the conceptual and external *schemas* (if needed) to meet *performance goals*.
- ❖ Start by understanding the *workload*:
 - Most important queries and how often they arise.
 - Most important updates and how often they arise.
 - Desired performance goals for those queries/updates?

Decisions to Be Made Include...



- ❖ What indexes should we create?
 - Which relations should have indexes? What field(s) should be their search keys? Should we build several indexes?
- ❖ For each index, what kind of an index should it be?
 - B+ tree? Hashed? Clustered? Unclustered?
- ❖ Should we make changes to the conceptual schema?
 - Consider alternative normalized schemas? (There are multiple choices when decomposing into BCNF, etc.)
 - Should we ``undo'' some decomposition steps and settle for a lower normal form? (*“Denormalization.”*)
 - Horizontal partitioning, materialized views, replication, ...

Understanding the Workload

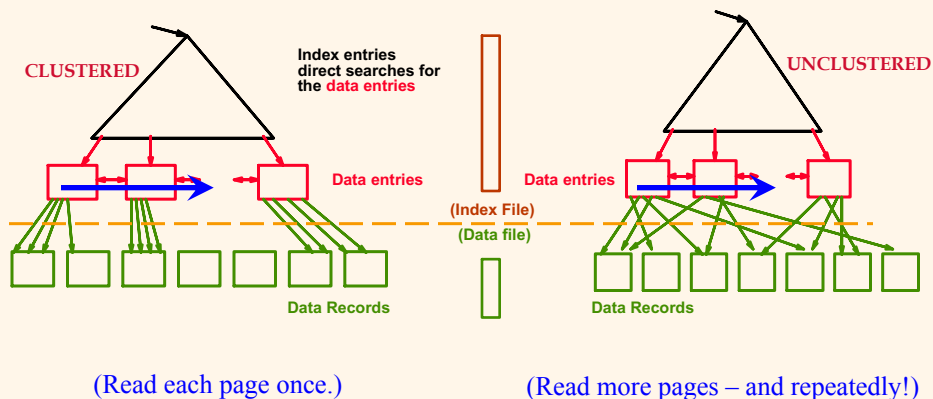


- ❖ For each query in the workload:
 - Which relations does it access?
 - Which attributes are retrieved?
 - Which attributes appear in selection/join conditions?
(And *how selective* are those conditions expected to be?)
- ❖ For each update in the workload:
 - Which attributes are involved in selection/join conditions?
(And *how selective* are those conditions likely to be?)
 - The type of update (INSERT/DELETE/UPDATE), and the attributes that are affected.


Index Classification (Review)

- ❖ *Primary vs. secondary*: If index search key contains the primary key, this is called the primary index.
 - *Unique* index: Search key contains a *candidate* key.
- ❖ *Clustered vs. unclustered*: If the order of data records is the same as, or 'close to', the order of stored data records, we have a clustered index.
 - A table can be clustered on *at most one* search key.
 - Cost of retrieving data records via an index varies *greatly* based on whether index is clustered or not!

Clustered vs. Unclustered Indexes (Reminder)




Choice of Indexes (Cont'd.)




- ❖ **One approach:** Consider the most important queries in turn. Consider the best query plan using the current indexes, and see if a better plan is possible with an additional index. If so, create it.
 - This means we must understand and see how a DBMS evaluates its queries. (*Query execution plans.*)
 - Let's start by discussing simple 1-table queries!
- ❖ Before creating an index, must also consider its impact on updates in the workload.
 - **Trade-off:** Indexes can make queries go faster, but updates will become slower. (Indexes require disk space, too.)

Index Selection Guidelines



- ❖ Attributes in **WHERE** clause are candidates for index keys.
 - Exact match condition → hashed index (or B+ tree if not).
 - Range query → B+ tree index.
 - Clustering especially useful for range queries, but can *also help with equality queries with duplicate values* (non-key field index).
- ❖ **Multi-attribute** search keys should be considered when a WHERE clause contains several conditions.
 - Order of attributes matters for range queries.
 - Such indexes can sometimes enable **index-only** strategies for important queries (e.g., aggregates / grouped aggregates).
 - *Note:* For index-only strategies, clustering isn't important!
- ❖ Choose indexes that benefit **as many queries** as possible.
 - Only **one** index can be clustered per relation, so choose it based on important queries that can benefit the most from clustering.

Examples of Clustered Indexes



```
SELECT E.dno
FROM   Emp E
WHERE  E.age > 40
```

- ❖ B+ tree index on E.age can be used to get qualifying tuples.
 - How selective is the condition?
 - Should the index be clustered?


```
SELECT E.dno,
       COUNT (*)
FROM   Emp E
WHERE  E.age > 10
GROUP BY E.dno
```

- ❖ Consider the GROUP BY query.
 - If most tuples have E.age > 10, using E.age index and sorting the retrieved tuples may be costly.
 - Clustered E.dno index may win!

```
SELECT E.dno
FROM   Emp E
WHERE  E.hobby='Stamps'
```

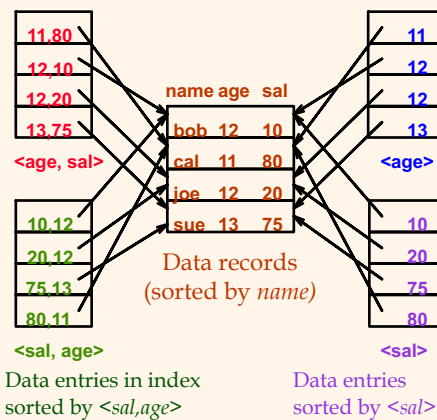
- ❖ Equality queries & duplicates:
 - Clustering on E.hobby helps!

Indexes with Composite Search Keys




- ❖ **Composite Search Keys:** Search on a combination of fields.
 - **Equality query:** Every field value is equal to a constant value. E.g. wrt <sal,age> index:
 - (age=20 AND sal=75)
 - **Range query:** Some field value is a range, not a constant. E.g. again wrt <sal,age> index:
 - age=20; or (age=20 AND sal > 10)
- ❖ Data entries in index sorted by search key to support such range queries.
 - **Lexicographic order**

Various composite key indexes using lexicographic (ASC) order.




Composite Search Keys



- ❖ To retrieve Emp records with $age=30$ AND $sal=4000$, an index on $\langle age, sal \rangle$ would be better than an index only on age or an index only on sal .
 - Note: Choice of index key is orthogonal to clustering.
- ❖ If condition is: $20 < age < 30$ AND $3000 < sal < 5000$:
 - Clustered B+ tree index on $\langle age, sal \rangle$ or $\langle sal, age \rangle$ is best.
- ❖ If condition is: $age=30$ AND $3000 < sal < 5000$:
 - Clustered $\langle age, sal \rangle$ index much better than $\langle sal, age \rangle$ index! (*Think about why!* Picture the index...)
- ❖ Composite indexes are larger; updated more often.

Index-Only Query Plans



- ❖ Some queries can be answered without retrieving *any* tuples from one or more of the relations involved if a suitable index is available.

(Sometimes called a "covering index" for the given query.)

$\langle E.dno \rangle$

```
SELECT E.dno, COUNT(*)
FROM Emp E
GROUP BY E.dno
```

$\langle E.dno, E.sal \rangle$

```
SELECT E.dno, MIN(E.sal)
FROM Emp E
GROUP BY E.dno
```

$\langle E.age, E.sal \rangle$



or

$\langle E.sal, E.age \rangle$

B+ tree index!

```
SELECT AVG(E.sal)
FROM Emp E
WHERE E.age=25 AND
      E.sal BETWEEN 3000 AND 5000
```

Some Illustrated Index-Only Plans



Note: The index files are each much smaller than the main file!