



Introduction to Data Management

Lecture #23 (Physical DB Design II)



Instructor: Mike Carey
mjcarey@ics.uci.edu



Announcements




❖ Two HW assignments remain:

- HW #7: Due tomorrow (5 PM)
 - Physical DB design (for MySQL and beyond)
- HW #8: Due next Thursday, June 6th (5 PM)!
 - NoSQL and **NoLateDay**

❖ Today's plan :

- Physical DB design wrap-up
- Then: **NoSQL & Big Data** (*a la* AsterixDB)
 - *Not* in the book,, so be sure to see the wiki for readings!
 - You can (and should) study ahead, e.g., by going through the Apache AsterixDB SQL++ Primer ([Using SQL++](#)).

Index Selection for Joins




❖ When considering a join condition:

- Index Nested Loop join (**INLJ**) method:
 - For each outer table tuple, use its join attribute value to probe the inner table for tuples to join (match) it with.
 - **Indexing the inner table's join column** will help!
 - Good for this index to be *clustered* if the join column is *not* the inner's PK (e.g., FK) and inner tuples need to be fetched.
- Sort-Merge join (**SMJ**) method:
 - Sort outer and inner tables on join attribute value and then scan them concurrently to match tuples.
 - **Clustered B+ trees on both join column(s)** fantastic for this!
- Hash join (**HJ**) method:
 - Indexing is not needed (not for the join, anyway).

Example 1

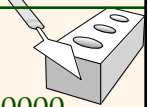
```
SELECT E.ename, D.mgr
FROM Emp E, Dept D
WHERE D.dname = 'Toy' AND E.dno=D.dno
```



- ❖ Hash index on *D.dname* supports 'Toy' selection.
 - Given this, an index on *D.dno* is not needed (not useful).
- ❖ Hash index on *E.dno* allows us to fetch matching (inner) Emp tuples for each outer Dept tuple.
- ❖ What if **WHERE** included: "... AND E.age=25"?
 - Could instead retrieve Emp tuples using index on *E.age*, then join with Dept tuples satisfying *dname* selection. (Comparable to strategy that uses the *E.dno* index.)
 - So, if *E.age* index were already created, this query provides less motivation for adding an *E.dno* index.

Example 2

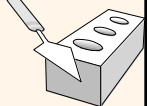
```
SELECT E.ename, D.mgr
FROM Emp E, Dept D
WHERE E.sal BETWEEN 10000 AND 20000
AND E.hobby= 'Stamps' AND E.dno=D.dno
```



- ❖ Clearly, Emp (E) should be the outer relation.
 - Suggests that we build an index (hashed) on *D.dno*.
- ❖ What index should we build on Emp?
 - B+ tree on *E.sal* could be used, OR an index on *E.hobby* could be used. Only one of these is needed, and which is better depends upon the *selectivity* of the conditions.
 - As a rough rule of thumb, equality selections tend to be more selective than range selections.
- ❖ As both examples indicate, our choice of indexes is guided by the plan(s) that we expect an optimizer to choose for a query. ∴ **Understand query optimizer!**


Clustering and Joins

```
SELECT E.ename, D.mgr
FROM Emp E, Dept D
WHERE D.dname= 'Toy' AND E.dno=D.dno
```




- ❖ Clustering is especially important when accessing inner tuples in INLJ (index nested loops join).
 - Should make index on *E.dno* clustered. (Q: See why?)
- ❖ Suppose that the WHERE clause were instead:
WHERE E.hobby= 'Stamps' AND E.dno=D.dno
 - If most employees collect stamps, Sort-Merge join may be worth considering. A *clustered* index on *D.dno* would help.
- ❖ **Summary:** Clustering is useful whenever *many* tuples are to be retrieved for one value or a range of values.

Tuning the Conceptual Schema



- ❖ The choice of conceptual schema should be guided by the workload, in addition to redundancy issues:
 - We may go for a 3NF (or lower!) schema rather than BCNF.
 - Workload may influence the choice we make in decomposing a relation into 3NF or BCNF.
 - We might *denormalize* (i.e., **undo** a decomposition step), or we might add fields to a relation.
 - We might consider *vertical decompositions*.
- ❖ If such changes come after a database is in use, it's called *schema evolution*; might want to mask some of the changes from applications by defining *views*.


Some Example Schemas (& Tradeoffs)



Suppliers(sid, sname, address, phone, ...)
Parts(pid, pname, size, color, listprice, ...)
Stock(sid, pid, price, quantity)


- ❖ What if a large fraction of the workload consists of Stock queries that also want suppliers' names?
 - `SELECT s.sid, s.sname, AVG(t.price) FROM Suppliers s, Stock t WHERE s.sid = t.sid GROUP BY s.sid, s.sname;`
 - Consider: `ALTER TABLE Stock ADD COLUMN sname ...;`
 - This is denormalization (on purpose, for performance!)
 - If $sid \rightarrow sname$ and $sname \rightarrow sid$, Stock would then be in 3NF.
 - Q: If $sid \rightarrow sname$ (but not vice versa), what NF would Stock be in?

Vertical Partitioning



- ❖ Consider a table with lots of columns, not all of which are of interest to all queries.
 - Ex: Emp(**eno**, email, name, addr, salary, age, dno)
- ❖ A given workload might actually turn out to be a “union of sub-workloads” in reality.
 - Employee communications queries
 - Employee compensation queries/analytics
 - Employee department queries/analytics

Vertical Partitioning Example



| eno | email | name | addr | salary | age | dno |
|-----|---------------|------|--------------|---------|-----|-----|
| 1 | joe@aol.com | Joe | 1 Main St. | 100000 | 25 | 10 |
| 2 | sue@gmail.com | Sue | 10 State St. | 125000 | 28 | 20 |
| 3 | zack@fb.com | Zack | 100 Wall St. | 2500000 | 40 | 30 |




(Vertical partitioning: ⋈)

| eno | email | name | addr | eno | salary | age | eno | dno |
|-----|---------------|------|--------------|-----|---------|-----|-----|-----|
| 1 | joe@aol.com | Joe | 1 Main St. | 1 | 100000 | 25 | 1 | 10 |
| 2 | sue@gmail.com | Sue | 10 State St. | 2 | 125000 | 28 | 2 | 20 |
| 3 | zack@fb.com | Zack | 100 Wall St. | 3 | 2500000 | 40 | 3 | 30 |


(In the limit: We get a column store!)

Horizontal Partitioning



- ❖ Occasionally, we may want to instead replace a relation by a set of relations that are *selections*.
 - Each new relation has same schema (columns) as the original, but only a subset of the rows.
 - Collectively, the new relations contain all rows of the original. (Typically, the new relations are *disjoint*.)
 - The original relation is the UNION (ALL) of the new ones (i.e., rather than the JOIN of the new ones).

Horizontal Partitioning Example



| eno | email | name | addr | salary | age | dno |
|-----|---------------|------|--------------|---------|-----|-----|
| 1 | joe@aol.com | Joe | 1 Main St. | 100000 | 25 | 10 |
| 2 | sue@gmail.com | Sue | 10 State St. | 125000 | 28 | 20 |
| 3 | zack@fb.com | Zack | 100 Wall St. | 2500000 | 40 | 30 |




(Horizontal partitioning: U)

| eno | email | name | addr | salary | age | dno |
|-----|---------------|------|--------------|--------|-----|-----|
| 1 | joe@aol.com | Joe | 1 Main St. | 100000 | 25 | 10 |
| 2 | sue@gmail.com | Sue | 10 State St. | 125000 | 28 | 20 |


| eno | email | name | addr | salary | age | dno |
|-----|-------------|------|--------------|---------|-----|-----|
| 3 | zack@fb.com | Zack | 100 Wall St. | 2500000 | 40 | 30 |

Potential Horizontal Rationale



- ❖ Suppose contracts with values over 10000 are subject to different rules. (This means queries on Contracts will frequently contain the condition *val* > 10000.)
- ❖ One approach to deal with this would be to create a clustered B+ tree index on Contracts(*val*).
- ❖ Another approach could be to replace Contracts by two relations, LargeContracts & SmallContracts, with the same attributes.
 - Performs like index but without index overhead.
 - Can then cluster on other (perhaps different!) attributes.

Masking Schema Changes




```
CREATE VIEW Contracts(cid, sid, jid, did, pid, qty, val)
AS SELECT * FROM LargeContracts
   UNION ALL
   SELECT * FROM SmallContracts
```

- ❖ Replacement of Contracts by LargeContracts and SmallContracts can be masked by this view.
- ❖ Note: queries with *val* > 10000 can be written against LargeContracts* for faster execution; users concerned with performance must be aware of this change.


(*The DBMS is unaware of the two tables' value constraints.)

In General: Tuning Queries (and Views)




- ❖ If a query runs slower than expected, see if an index needs to be re-built, or if **table statistics** are too old.
- ❖ Sometimes, the DBMS may not be executing the plan you had in mind. Common areas of weakness:
 - Selections involving **arithmetic** or **LIKE expressions**.
 - Selections involving **OR** conditions.
 - Selections involving **null values**.
 - **Lack of advanced evaluation features** like some index-only strategies or certain join methods, or poor size estimation.
- ❖ Check the query plan!!! Then adjust the choice of indexes or maybe **rewrite the query or view**.

Miscellany for Query Tuning




- ❖ Minimize the use of **DISTINCT**: Don't use the D-word if duplicates are acceptable or if the answer contains a key.
- ❖ Consider the DBMS's use of indexes when writing arithmetic expressions: $E.age = 2 * D.age$ will benefit from an index on **E.age**, but it probably wouldn't benefit from an index on **D.age**!

Physical DB Design Summary




- ❖ End-to-end DB design consists of several tasks: *requirements analysis, conceptual design, schema refinement, physical design and finally tuning.*
 - In general, one goes back and forth between tasks to refine a DB design
 - Decisions made in one task can influence choices in another task.
- ❖ Understanding the *workload* for the application, and performance goals, is essential to good design.
 - What are the important queries and updates?
What attributes/relations are involved?

Summary (Cont'd.)



- ❖ The conceptual schema should perhaps be refined by considering performance criteria and workload:
 - May choose 3NF or a lower normal form over BCNF.
 - May choose among several alternative decompositions based on the expected workload.
 - May actually *denormalize*, or *undo*, some decompositions.
 - May consider further *vertical* or *horizontal* decompositions.

Summary (Cont'd.)



- ❖ Over time, indexes may have to be fine-tuned (dropped, created, re-built, ...) for performance.
 - Be sure to examine the query plan(s) used by the system and adjust the choices of indexes appropriately.
- ❖ Sometimes the system may still not find a good plan:
 - Null values, arithmetic conditions, string expressions, the use of ORs, etc., can “confuse” some query optimizers.
- ❖ So, may have to rewrite a particular query or view:
 - Might need to re-examine your complex nested queries, complex conditions, or operations like DISTINCT.
- ❖ Any lingering questions...?