


Introduction to Data Management

Lecture #24

~~SQL~~ **NoSQL**



Instructor: Mike Carey

mjcarey@ics.uci.edu

1

It's time for the season finale of...

Friday Nights With
Databases...

Brought to you by...



2

Announcements

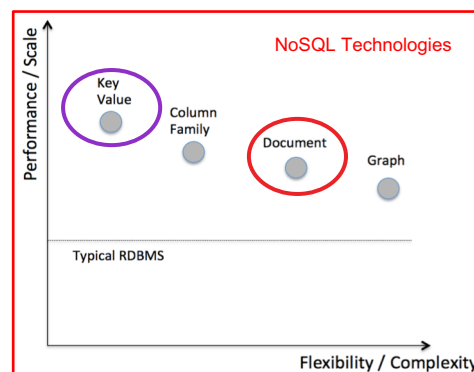
- Homework info:
 - HW #8 (NoSQL) is the **last one!**
 - Due Thursday (at 5 PM), with **NoLateDay** (!)
 - Warning: *LOAD* can be (path) finicky...
- Endterm exam info:
 - Non-cumulative and **in class next Friday** (as usual)
- NoSQL lecture plans:
 - Today: **NoSQL & Big Data** (*a la* AsterixDB)
 - Refer to the **Using SQL++** Primer and other docs on the Apache AsterixDB site
 - Read **SQL++ For SQL Users** from Couchbase, by **Don Chamberlin** (the Father of SQL!)
 - **Lots** of useful info for moving from SQL to SQL++!



3

What is a **NoSQL DB** – why “not SQL”?

- **Not** from the DB world
 - Distributed systems folks
 - Also various startup companies
- From caches → persistent K/V use cases
 - Apps needed massive scale-out
 - OLTP (vs. parallel query DB) apps
 - Simple, low-latency API – get/put by key
 - Need a key K, but want **no schema** for V
 - Record-level atomicity, replica consistency varies
- In the context of this talk, NoSQL does **not** mean
 - Hadoop (or SQL on Hadoop)
 - Graph databases or graph analytics platforms



4

NoSQL Data (JSON-based)

Collection(Orders)

```
{ "id": "123",
  "Customer":
    { "custName": "Fred",
      "custCity": "LA" }
  "total": 25.97,
  "Items": [
    { "product-sku": 401,
      "qty": 2,
      "price": 9.99 },
    { "product-sku": 544,
      "qty": 1,
      "price": 3.99 }
  ]
}
```

Collection(Products)

```
{ "sku": 401,
  "name": "Garfield T-Shirt",
  "listPrice": 9.99,
  "size": "XL" },
{ "sku": 544,
  "name": "USB Charger",
  "listPrice": 5.99,
  "power": "115V" }
```


5

Current NoSQL (document DB) trends

- Popular examples: MongoDB, Couchbase
- Users now coveting the benefits of many DB goodies
 - Secondary indexing and non-key access
 - Declarative queries
 - Aggregates and now (commonly small) joins
- World seems to be heading towards...
 - BDMS (think scalable, OLTP-aimed, parallel/distributed DBMS)
 - Declarative queries and query optimization, applied to schema-less data
 - Return of (some, optional!) schema information

6

Towards a Big Data Management System (BDMS)

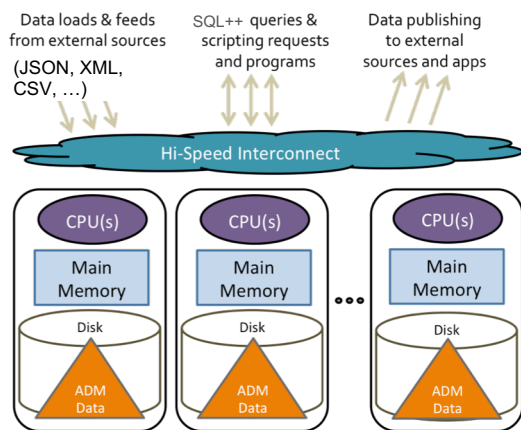


BDMS Desiderata:

- Able to **manage** data
- **Flexible** data model
- Full **query** capability
- Continuous data **ingestion**
- Efficient and robust **parallel** runtime
- Cost **proportional** to task at hand
- Support "**Big Data** data types"
-
-
-

7

Apache AsterixDB (from UCI+UCR)



ASTERIX Goal:
To ingest, digest,
persist, index,
manage, query,
analyze, and
publish massive
quantities of
semistructured
information...



<http://asterixdb.apache.org/>



8

Data Model: JSON (JavaScript Object Notation)

Customers

```
{
  "custid": "C37",
  "name": "T. Hanks",
  "address": {
    "street": "120 Harbor Blvd.",
    "city": "Boston, MA",
    "zipcode": "02115"
  },
  "rating": 750
}

{
  "custid": "C47",
  "name": "S. Lauren",
  "address": {
    "street": "17 Rue d'Antibes",
    "city": "Cannes, France"
  },
  "rating": 625
}
```

Orders

```
{
  "orderno": 1004,
  "custid": "C35",
  "order_date": "2017-07-10",
  "ship_date": "2017-07-15",
  "items": [
    {
      "itemno": 680,
      "qty": 6,
      "price": 9.99
    },
    {
      "itemno": 195,
      "qty": 4,
      "price": 35.00
    }
  ]
}
```

...

```
{
  "orderno": 1008,
  "custid": "C13",
  "order_date": "2017-10-13",
  "items": [
    {
      "itemno": 460,
      "qty": 20,
      "price": 99.99
    }
  ]
}
```

Data from D. Chamberlin. SQL++ for SQL Users: A Tutorial

9

Data

Customers

```
{
  "custid": "C37",
  "name": "T. Hanks",
  "address": {
    "street": "120 Harbor Blvd.",
    "city": "Boston, MA",
    "zipcode": "02115"
  },
  "rating": 750
}

{
  "custid": "C47",
  "name": "S. Lauren",
  "address": {
    "street": "17 Rue d'Antibes",
    "city": "Cannes, France"
  },
  "rating": 625
}
```

Orders

```
{
  "orderno": 1004,
  "custid": "C35",
  "order_date": "2017-07-10",
  "ship_date": "2017-07-15",
  "items": [
    {
      "itemno": 680,
      "qty": 6,
      "price": 9.99
    },
    {
      "itemno": 195,
      "qty": 4,
      "price": 35.00
    }
  ]
}
```

...

```
{
  "orderno": 1008,
  "custid": "C13",
  "order_date": "2017-10-13",
  "items": [
    {
      "itemno": 460,
      "qty": 20,
      "price": 99.99
    }
  ]
}
```

Data from D. Chamberlin. SQL++ for SQL Users: A Tutorial

10

Data (Relational version)

Customers

```
{
  "custid": "C37",
  "name": "T. Hanks",
  "address_street": "120 Harbor Blvd.",
  "address_city": "Boston, MA",
  "address_zipcode": "02115",
  "rating": 750
}
{
  "custid": "C47",
  "name": "S. Lauren",
  "address_street": "17 Rue d'Antibes",
  "address_city": "Cannes, France",
  "address_zipcode": null,
  "rating": 625
}
```

Orders

```
{
  "orderno": 1004,
  "custid": "C35",
  "order_date": "2017-07-10",
  "ship_date": "2017-07-15"
}
{
  "orderno": 1008,
  "custid": "C13",
  "order_date": "2017-10-13",
  "ship_date": null
}
```

```
CREATE TABLE Lineitems (
  orderno INTEGER,
  itemno INTEGER,
  quantity INTEGER NOT NULL,
  price DECIMAL(8,2) NOT NULL,
  PRIMARY KEY (orderno, itemno),
  FOREIGN KEY (orderno) REFERENCES Orders(orderno)
)
```

Lineitems

```
{
  "orderno": 1004,
  "itemno": 680,
  "qty": 6,
  "price": 9.99
}
{
  "orderno": 1004,
  "itemno": 195,
  "qty": 4,
  "price": 35.00
}
{
  "orderno": 1008,
  "itemno": 460,
  "qty": 20,
  "price": 99.99
}
```

11

Data (Relational version)

Customers

```
{
  "custid": "C37",
  "name": "T. Hanks",
  "address_street": "120 Harbor Blvd.",
  "address_city": "Boston, MA",
  "address_zipcode": "02115",
  "rating": 750
}
{
  "custid": "C47",
  "name": "S. Lauren",
  "address_street": "17 Rue d'Antibes",
  "address_city": "Cannes, France",
  "address_zipcode": null,
  "rating": 625
}
```

Orders

```
{
  "orderno": 1004,
  "custid": "C35",
  "order_date": "2017-07-10",
  "ship_date": "2017-07-15"
}
{
  "orderno": 1008,
  "custid": "C13",
  "order_date": "2017-10-13",
  "ship_date": null
}
```

```
CREATE TABLE Lineitems (
  orderno INTEGER,
  itemno INTEGER,
  quantity INTEGER NOT NULL,
  price DECIMAL(8,2) NOT NULL,
  PRIMARY KEY (orderno, itemno),
  FOREIGN KEY (orderno) REFERENCES Orders(orderno)
)
```

Lineitems

```
{
  "orderno": 1004,
  "itemno": 680,
  "qty": 6,
  "price": 9.99,
  "currency": "USD"
}
{
  "orderno": 1004,
  "itemno": 195,
  "qty": 4,
  "price": 35.00,
  "currency": "USD"
}
{
  "orderno": 1008,
  "itemno": 460,
  "qty": 20,
  "price": 99.99,
  "currency": "EUR"
}
```

12

Sloppy Data

Customers

```
{
  "custid": "C37",
  "name": "T. Hanks",
  "address": {
    "street": "120 Harbor Blvd.",
    "city": "Boston, MA",
    "zipcode": "02115"
  },
  "rating": 750
}

{
  "custid": "C47",
  "name": "S. Lauren",
  "address": {
    "street": "17 Rue d'Antibes",
    "city": "Cannes, France"
  },
  "rating": "625"
}
```

Orders

```
{
  "orderno": 1004,
  "custid": "C35",
  "order_date": "2017-07-10",
  "ship_date": "2017-07-15",
  "items": [
    {
      "itemno": 680,
      "qty": 6,
      "price": 9.99
    },
    {
      "itemno": 195,
      "qty": 4,
      "price": "if you have to ask ..."
    }
  ]
}
```

...

```
{
  "orderno": 1008,
  "custid": "C13",
  "order_date": "2017-10-13",
  "items": {
    "itemno": 460,
    "qty": 20,
    "price": 99.99
  }
}
```

13

SQL++: Just like SQL ...

```
SELECT name
FROM customers
WHERE rating > 650;
```

```
[
  {
    "name": "M. Streep"
  },
  {
    "name": "T. Hanks"
  },
  {
    "name": "T. Cruise"
  }
]
```

14

Just like SQL ...

```
SELECT name
FROM customers
WHERE rating > 650;
```

```
SELECT c.name, o.order_date
FROM customers AS c, orders AS o
WHERE c.custid = o.custid
AND c.custid = "C41";
```

```
[
  {
    "name": "R. Duvall",
    "order_date": "2017-09-02"
  },
  {
    "name": "R. Duvall",
    "order_date": "2017-04-29"
  }
]
```

15

Just like SQL ...

```
SELECT name
FROM customers
WHERE rating > 650;
```

```
SELECT c.name, o.order_date
FROM customers AS c, orders AS o
WHERE c.custid = o.custid
AND c.custid = "C41";
```

```
SELECT c.name, o.order_date
FROM customers AS c LEFT OUTER JOIN orders AS o
ON c.custid = o.custid
WHERE c.custid = "C41";
```

16

Just like SQL ...

```
SELECT name
FROM customers
WHERE rating > 650;
```

```
SELECT c.name, o.order_date
FROM customers AS c, orders AS o
WHERE c.custid = o.custid
      AND c.custid = "C41";
```

```
SELECT order_date, count(*) AS cnt
FROM orders
GROUP BY order_date
HAVING count(*) > 0
ORDER BY order_date DESC
LIMIT 3;
```

```
[
  {
    "cnt": 1,
    "order_date": "2017-10-13"
  },
  {
    "cnt": 1,
    "order_date": "2017-09-13"
  },
  {
    "cnt": 1,
    "order_date": "2017-09-02"
  }
]
```

17

... almost!

```
SELECT name, order_date
FROM customers, orders
WHERE customers.custid = orders.custid
      AND rating > 650;
```

Cannot resolve ambiguous alias reference for
identifier rating (in line 4, at column 7)

18

... almost!

```
SELECT name, order_date
FROM customers, orders
WHERE customers.custid = orders.custid
    AND rating > 650;
```

```
SELECT c.name, o.order_date
FROM customers AS c, orders AS o
WHERE c.custid = o.custid
    AND c.rating > 650;
```

```
[
  {
    "name": "T. Hanks",
    "order_date": "2017-08-30"
  },
  {
    "name": "T. Cruise",
    "order_date": "2017-05-01"
  },
  {
    "name": "T. Cruise",
    "order_date": "2017-10-13"
  },
  {
    "name": "T. Cruise",
    "order_date": "2017-09-13"
  }
]
```

19

... almost!

```
SELECT name, order_date
FROM customers, orders
WHERE customers.custid = orders.custid
    AND rating > 650;
```

```
SELECT c.name, o.order_date
FROM customers AS c, orders AS o
WHERE c.custid = o.custid
    AND c.rating > 650;
```

```
SELECT *
FROM customers AS c, orders AS o
WHERE c.custid = o.custid
    AND c.rating > 650;
```

```
[
  {
    "c": {
      "address": {
        "city": "Boston, MA",
        "street": "120 Harbor Blvd.",
        "zipcode": "02115"
      },
      "custid": "C37",
      "name": "T. Hanks",
      "rating": 750
    },
    "o": {
      "custid": "C37",
      "items": [
        {
          "itemno": 460,
          "price": 99.98,
          "qty": 2
        }
      ]
    }
  },
  ...
]
```

20

SELECT VALUE: Added "VALUE"

```
SELECT VALUE name
FROM customers
WHERE rating > 650;
```

```
[
  "M. Streep",
  "T. Hanks",
  "T. Cruise"
]
```

21

Added "VALUE"

```
SELECT VALUE name
FROM customers
WHERE rating > 650;
```

```
SELECT VALUE {
  "CustomerName":c.name,
  "OrderDate":o.order_date
}
FROM customers AS c, orders AS o
WHERE c.custid = o.custid
AND c.rating > 650;
```

```
[
  {
    "CustomerName": "T. Hanks",
    "OrderDate": "2017-08-30"
  },
  {
    "CustomerName": "T. Cruise",
    "OrderDate": "2017-09-13"
  },
  {
    "CustomerName": "T. Cruise",
    "OrderDate": "2017-05-01"
  },
  {
    "CustomerName": "T. Cruise",
    "OrderDate": "2017-10-13"
  }
]
```

22

Added "VALUE"

```
SELECT VALUE name
FROM customers
WHERE rating > 650;
```

```
SELECT VALUE {
  "CustomerName":c.name,
  "OrderDate":o.order_date
}
FROM customers AS c, orders AS o
WHERE c.custid = o.custid
AND c.rating > 650;
```

```
SELECT c.name AS CustomerName,
       o.order_date AS OrderDate
FROM customers AS c, orders AS o
WHERE c.custid = o.custid
AND c.rating > 650;
```

23

Added "VALUE"

```
SELECT VALUE name
FROM customers
WHERE rating > 650;
```

```
SELECT VALUE {
  "CustomerName":c.name,
  "OrderDate":o.order_date
}
FROM customers AS c, orders AS o
WHERE c.custid = o.custid
AND c.rating > 650;
```

```
SELECT VALUE {
  "CustomerName":c.name,
  "Orders":(SELECT VALUE o.orderno FROM orders AS o
            WHERE o.custid = c.custid)
}
FROM customers AS c
WHERE c.custid = "C41";
```

```
[
  {
    "Orders": [
      1006,
      1001
    ],
    "CustomerName": "R. Duvall"
  }
]
```

24

Quiz

Which query retrieves the names of the customers that have the highest rating?

- A
- ```
SELECT name
FROM customers
WHERE rating =
 (SELECT MAX(rating) FROM customers);
```
- B
- ```
SELECT c1.name
FROM customers AS c1
WHERE c1.rating =
  (SELECT VALUE MAX(c2.rating) FROM customers AS c2);
```
- C
- ```
SELECT c1.name
FROM customers AS c1
WHERE c1.rating =
 (SELECT MAX(c2.rating) FROM customers AS c2);
```
- D
- ```
SELECT VALUE c1.name
FROM customers AS c1
WHERE c1.rating =
  (SELECT VALUE MAX(c2.rating) FROM customers AS c2)[0];
```

25

SQL Pitfalls and the value of VALUE

```
SELECT name
FROM customers
WHERE rating =
  (SELECT MAX(rating) FROM customers);
```

Type mismatch: expected value of type multiset or array, but got the value of type object (in line 4, at column 28)

26

SQL Pitfalls and the value of VALUE

```
SELECT name
FROM customers AS c
WHERE rating =
  (SELECT MAX(rating) FROM c);
```

Type mismatch: expected value of type multiset or array, but got the value of type object (in line 4, at column 28)

27

SQL Pitfalls and the value of VALUE

```
SELECT name
FROM customers
WHERE rating =
  (SELECT MAX(rating) FROM customers);
```

[]

```
SELECT c1.name
FROM customers AS c1
WHERE c1.rating =
  (SELECT MAX(c2.rating) FROM customers AS c2);
```

28

SQL Pitfalls and the value of VALUE

```

SELECT name                                [ ]
FROM customers
WHERE rating =
    (SELECT MAX(rating) FROM customers);

SELECT c1.name
FROM customers AS c1
WHERE c1.rating =
    (SELECT MAX(c2.rating) FROM customers AS c2);

SELECT c1.name
FROM customers AS c1
WHERE c1.rating =
    (SELECT VALUE MAX(c2.rating) FROM customers AS c2);

```

29

SQL Pitfalls and the value of VALUE

```

SELECT name                                [
FROM customers                             "T. Cruise",
WHERE rating =                             "T. Hanks"
    (SELECT MAX(rating) FROM customers);    ]

SELECT c1.name
FROM customers AS c1
WHERE c1.rating =
    (SELECT MAX(c2.rating) FROM customers AS c2);

SELECT c1.name
FROM customers AS c1
WHERE c1.rating =
    (SELECT VALUE MAX(c2.rating) FROM customers AS c2);

SELECT VALUE c1.name
FROM customers AS c1
WHERE c1.rating =
    (SELECT VALUE MAX(c2.rating) FROM customers AS c2)[0];

```

30

Quiz

Which query retrieves the names of the customers that have the highest rating?

- A
- ```
SELECT name
FROM customers
WHERE rating =
 (SELECT MAX(rating) FROM customers);
```
- B
- ```
SELECT c1.name
FROM customers AS c1
WHERE c1.rating =
  (SELECT VALUE MAX(c2.rating) FROM customers AS c2);
```
- C
- ```
SELECT c1.name
FROM customers AS c1
WHERE c1.rating =
 (SELECT MAX(c2.rating) FROM customers AS c2);
```
- D**
- ```
SELECT VALUE c1.name
FROM customers AS c1
WHERE c1.rating =
  (SELECT VALUE MAX(c2.rating) FROM customers AS c2)[0];
```

31

To be continued....

32

More information about JSON, SQL++, and AsterixDB

- Asterix project UCI/UCR research home
 - <http://asterix.ics.uci.edu/>
- Apache AsterixDB home
 - <http://asterixdb.apache.org/>
- SQL++ Primer
 - <https://ci.apache.org/projects/asterixdb/sqlpp/primer-sqlpp.html>
- Navigate from CS122a wiki (HW) to get and install it...!
 - Also, a few other resources and hints in the HW materials

